

Kubernetes Object and Service

SoC Summer Workshop
Cloud Computing with Big Data

Richard T. B. Ma

School of Computing

National University of Singapore

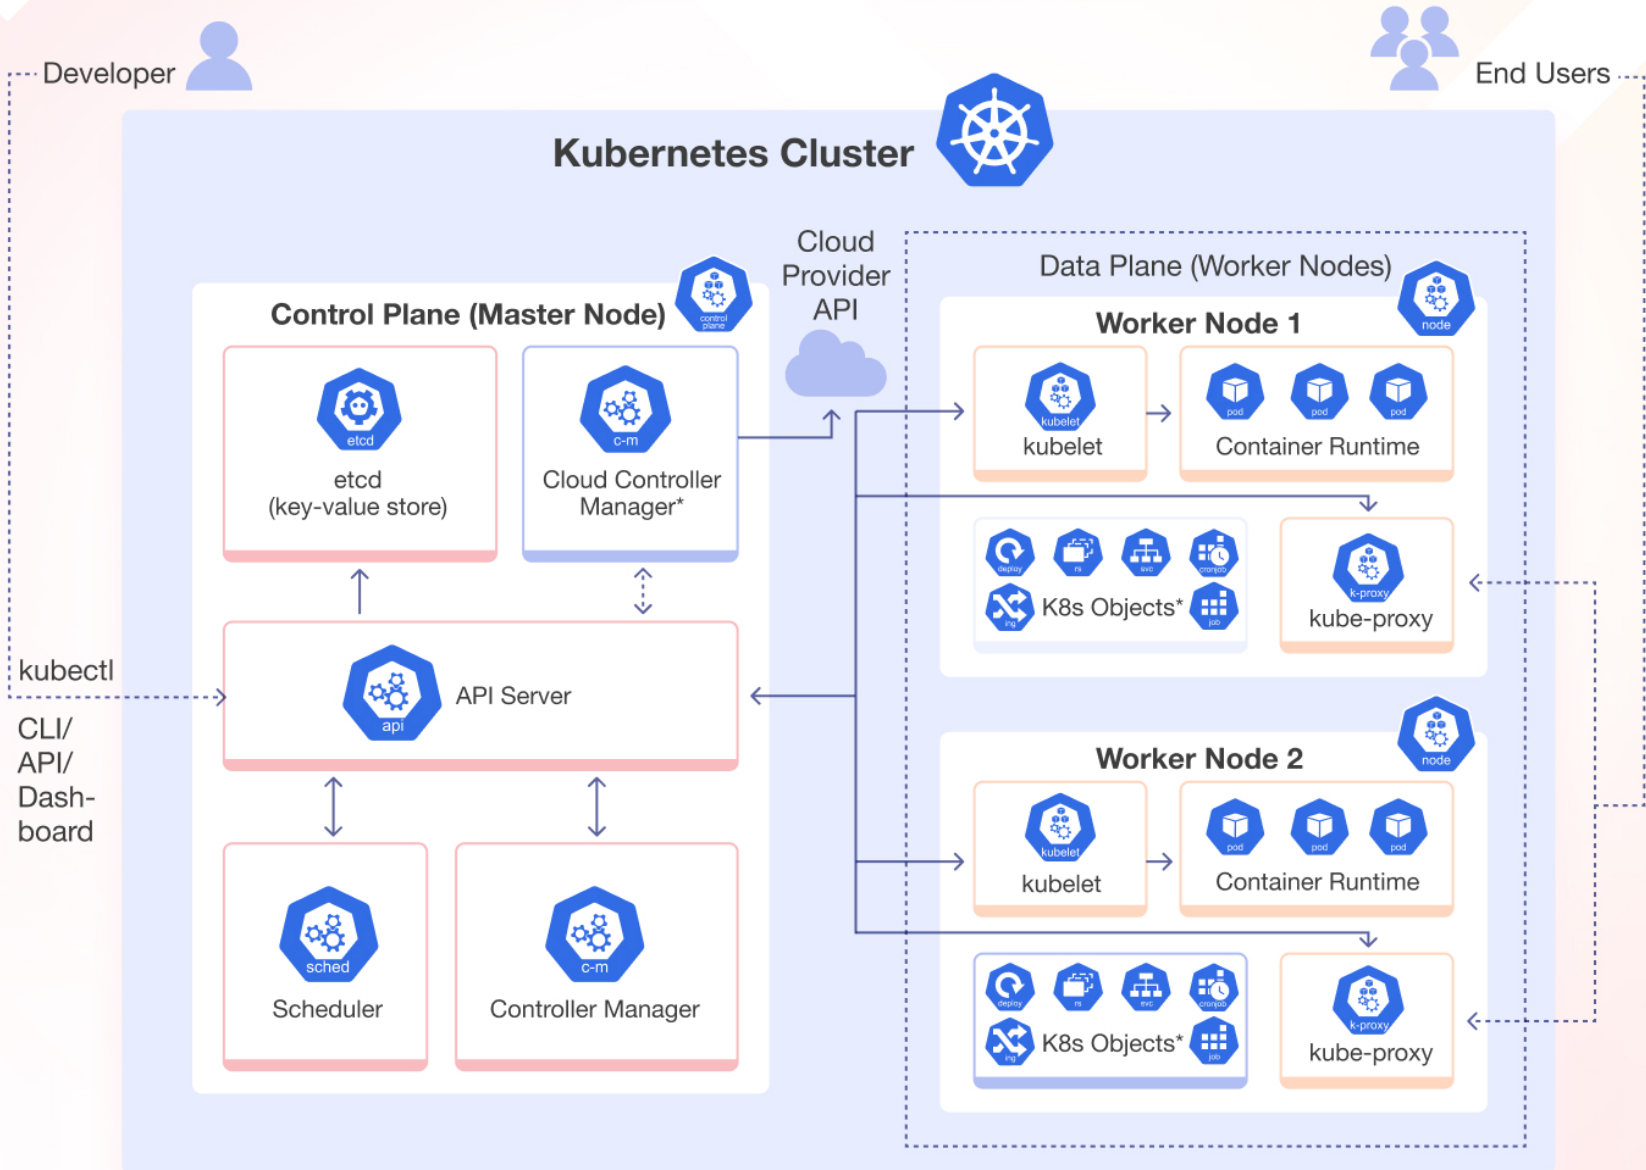
Formation of Project Groups

- ❑ Each project group consists of 4 students
 - ❖ with two groups that have 5 students
 - ❖ to be fair, each third-year student will lead a group with second-year students
- ❑ Potential topics of projects
 - ❖ traditional applications
 - ❖ new applications over the cloud
- ❑ Types of projects
 - ❖ cloud system with orchestration
 - ❖ cloud-native big data application

Roadmap

- ❑ Quick Review
- ❑ Kubernetes Object and Manifest Files
- ❑ Kubernetes Service
- ❑ We'll have a tutorial tomorrow on
 - ❖ Amazon and Google Container Engine (GKE)

Kubernetes Architecture



*optional

Kubernetes Node Components

1. **Kubelet**: an agent responsible for managing the node, communicating with the API server, and ensuring that containers are running as specified.
2. **Container Runtime**: a container runtime, such as Docker, containerd, or CRI-O, is responsible for pulling container images and running containers.
3. **Kube-proxy**: a network proxy responsible for implementing K8s service abstraction by maintaining network rules and forwarding traffic to the appropriate pods.
4. **Pod CIDR**: a range of IP addresses assigned, can be allocated to pods running on that node.

Kubernetes Node Operations

1. **Pod Scheduling:** Nodes are responsible for hosting and running pods. When a pod is scheduled to run, the K8s scheduler selects a node based on resource requirements, affinity/anti-affinity rules, node selectors, and other constraints.
2. **Status and Health:** Nodes report their status and health to the API server, e.g., resource utilization, node conditions (e.g., Ready, NotReady), and any issues or errors encountered on the node.
3. **Scaling and Capacity:** K8s clusters can dynamically add or remove nodes to accommodate changing workload demands, i.e., horizontal scaling.


Roadmap

- ❑ Quick Review
- ❑ Kubernetes Object and Manifest Files
- ❑ Kubernetes Service


<p>Deployment Unit</p>  <p>Pod</p>	<p>Networking</p>  <p>Service</p>	<p>Storage</p>  <p>Volume</p>	<p>Isolation</p>  <p>Namespace</p>
<p>Replication</p>  <p>Replicaset</p>	<p>POD Management</p>  <p>Deployment</p>	<p>State Management</p>  <p>StatefulSet</p>	<p>Node Operation</p>  <p>DaemonSet</p>
<p>Task Execution</p>  <p>Job</p>	<p>Scheduled Tasks</p>  <p>CronJob</p>	<p>Sensitive Data</p>  <p>Secret</p>	<p>Configuration</p>  <p>ConfigMap</p>
<p>External Access</p>  <p>Ingress</p>	<p>Network Rules</p>  <p>NetworkPolicy</p>	<p>Persistent Storage</p>  <p>Persistent Volume</p>	<p>Storage Request</p>  <p>PV Claim</p>
<p>Scalability</p>  <p>HPA/VPA</p>	<p>Network endpoints</p>  <p>EndpointSlices</p>	<p>Authentication</p>  <p>ServiceAccount</p>	<p>Authorization</p> <p>Role/ClusterRole</p>

Objects and Controllers


Kubernetes
objects



update



watch




watch



update

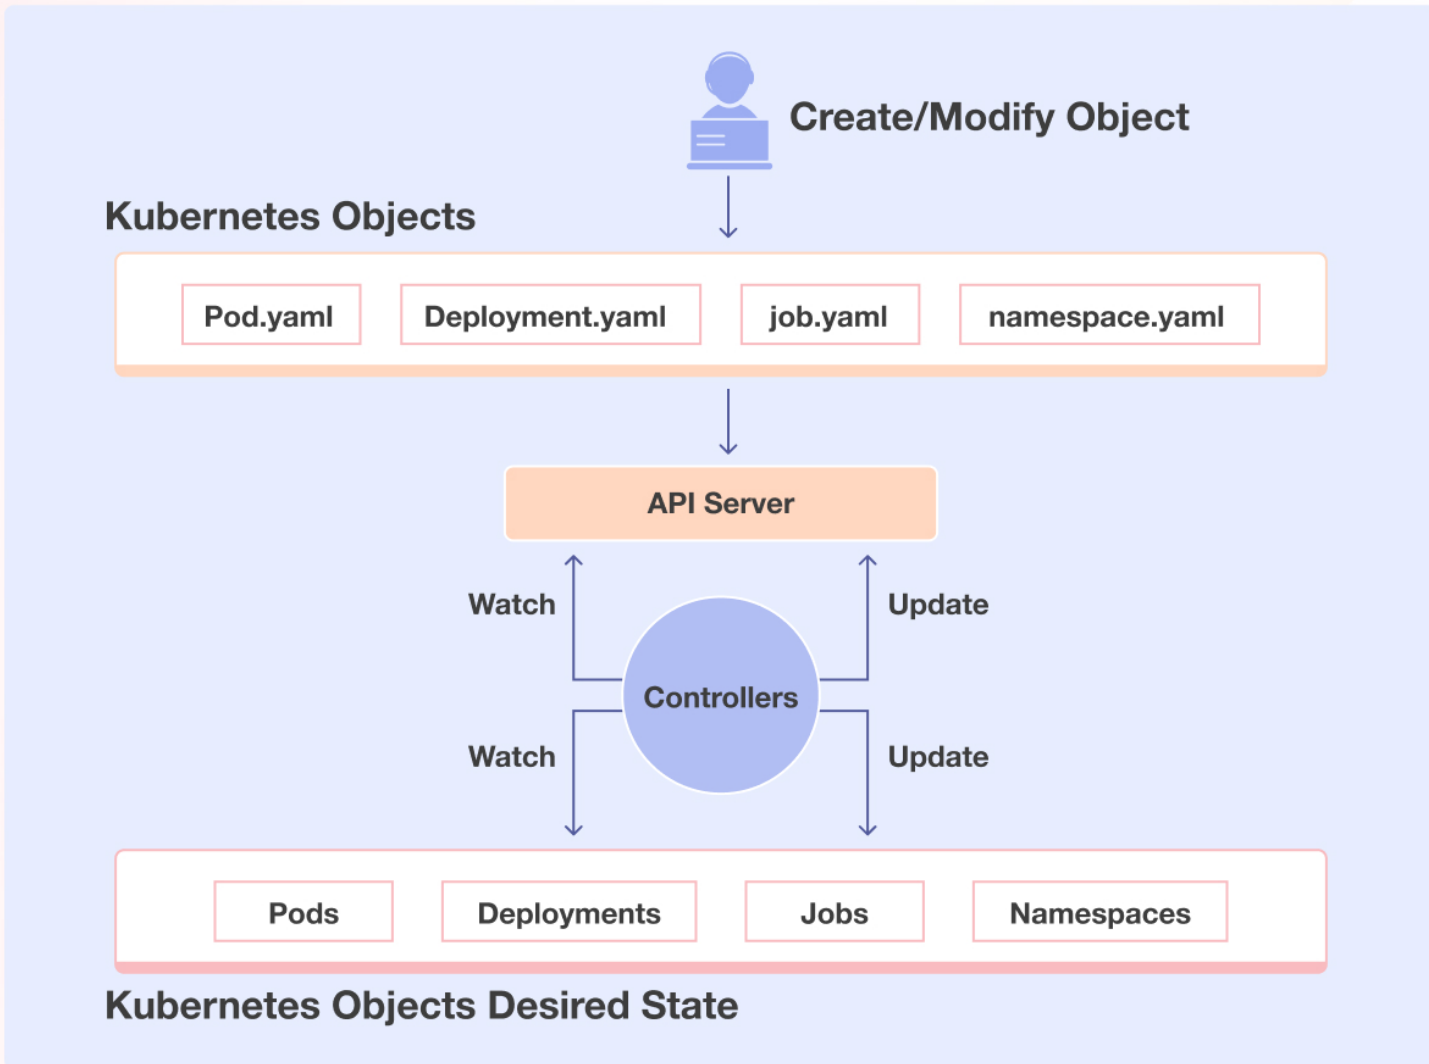
System
Resources



Kubernetes Object

- ❑ Persistent entities in K8s, representing a "record of intent" that tells what you want
 - ❖ described as a "specification" or "spec"
 - ❖ K8s manages to reach the *desired state*.
- ❑ The *current* state, i.e., "status", of cluster
 - ❖ what apps are running on which nodes
 - ❖ the resources available to those apps
 - ❖ the policies around how the apps behave, e.g., restart policies, upgrades, and fault-tolerance
- ❑ Use K8s API via `kubectl` or client libs.

kube-controller-manager



Structure of YAML Manifest Files



```
apiVersion: v1
kind: Pod
metadata:
  name: server-pod-v1
spec:
  containers:
  - name: server-container
    image: yancanmao/server-image
```

❑ apiVersion

❑ kind

❖ type of object

❑ metadata

❖ name, label, & etc.

❑ spec

❖ kind-specific defs

❑ status

❖ runtime states

```
$ kubectl get pod server-pod-v1 -o yaml
```

```

apiVersion: v1
kind: Pod
metadata:
  name: server-pod-v1
spec:
  containers:
  - image: yancanmao/server-image
    name: server-container
status:
  conditions:
  - lastProbeTime: null
    lastTransitionTime: "2024-05-19T03:07:31Z"
    status: "True"
    type: PodReadyToStartContainers
  - lastProbeTime: null
    lastTransitionTime: "2024-05-19T03:07:28Z"
    status: "True"
    type: Initialized
  - lastProbeTime: null
    lastTransitionTime: "2024-05-19T03:07:31Z"
    status: "True"
    type: Ready
  - lastProbeTime: null
    lastTransitionTime: "2024-05-19T03:07:31Z"
    status: "True"
    type: ContainersReady
  - lastProbeTime: null
    lastTransitionTime: "2024-05-19T03:07:28Z"
    status: "True"
    type: PodScheduled
  containerStatuses:
  - containerID: docker://57789f43840a65ae7e3290dc023230c381f2fc6146516b13ae366d96acbb700a
    image: yancanmao/server-image:latest
    imageID: docker-pullable://yancanmao/server-
image@sha256:1d79b0f25e226f852ccdae1b65af4a069e77699af3142b53cc084316196b7d3c
    lastState: {}
    name: server-container
    ready: true
    restartCount: 0
    started: true
    state:
      running:
        startedAt: "2024-05-19T03:07:30Z"
  hostIP: 192.168.49.2
  hostIPs:
  - ip: 192.168.49.2
  phase: Running
  podIP: 10.244.0.10
  podIPs:
  - ip: 10.244.0.10
  qosClass: BestEffort
  startTime: "2024-05-19T03:07:28Z"

```

Object Status

❑ apiVersion

❑ kind

❖ object & non-object

❑ metadata

❖ name, label, & etc.

❑ spec

❖ kind-specific defs

❑ status

❖ runtime states

Object Metadata

❑ Every object kind **MUST** have:

- ❖ **name**: a string that uniquely identifies this object within the current namespace, used in the path when retrieving an object.
- ❖ **namespace**: a DNS-like label that objects are subdivided into.
- ❖ **uid**: a unique value in *time* & space used to distinguish between objects of the same name that have been deleted & recreated.

❑ Every object kind **SHOULD** have:

- ❖ **labels**: a map of string keys and values, i.e., key-value pairs, that can be used to organize and categorize objects.
- ❖ **annotations**: key-value pairs, can be used by external tooling to store and retrieve arbitrary metadata about this object.
- ❖ **resourceVersion**: the internal version of this object, used to guarantee consistent updating.

Namespace

- ❑ A mechanism for isolating groups of resources within a cluster.
 - ❖ Quotas can be defined for each namespace to limit the resources consumed.
 - ❖ Unique names for resources within a namespace, but not across namespaces
- ❑ Scoping is applicable for namespaced objects, e.g., Service, not for cluster-wide objects, e.g. Node.
 - ❖ Resources within the namespaces can refer to each other with their service names.
 - ❖ Resources across namespace can be reached via:
`<service_name>.<namespace_name>.svc.cluster.local`

Namespace - how to use?

- ❑ Create a namespace:

```
$ kubectl create namespace <namespace_name>
```

- ❑ The active namespace is recorded in the "context"

- ❖ a configuration file for kubectl

- ❖ located at \$HOME/.kube/config

- ❖ also include info about the cluster and users

- ❑ Change the namespace in the current context:

```
$ kubectl config set-context --current --  
namespace=<namespace_name>
```

- ❑ Change context:

```
$ kubectl config use-context <context_name>
```


Object Spec ↔ Object Status

- ❑ Spec: the desired state of an object
 - ❖ configs from users
 - ❖ default values by system
 - ❖ initialize/updated by other components, e.g., scheduler, auto-scaler
 - persisted in storage with the API object
 - ❖ if deleted, the object will be purged from the system
- ❑ Status: current state of the object in system
 - ❖ usually persisted with the object by automated processes, generated on the fly
 - ❖ should be the most recent observations
 - ❖ may contain info such as the results of operations executed in response to the spec

Roadmap

- ❑ Quick Review
- ❑ Kubernetes Object and Manifest Files
- ❑ Kubernetes Service

How to interact with the containers?

- ❑ Remote login to execute and interact with programs inside containers

```
$ kubectl exec mypod -stdin -tty -- /bin/bash
```

- ❑ Access from another pod within the cluster
 - ❖ need to know the internal IP of the destination pod

- ❑ What if I want to access from my local machine?
 - ❖ port forwarding: access a specific port of a container running inside a Kubernetes Pod from your local machine


```
$ kubectl port-forward --address 0.0.0.0 mypod 8888:8080
```

- ❖ for temporary access during development or debugging


Service Discovery

❑ Traditional client-side service discovery

- ❖ an agent capable of looking at a registry for service instances and choosing one to call



❑ K8s: server-side service discovery



- ❖ consumer calls a fixed virtual Service endpoint that can dynamically discover service instances

Kubernetes Service

- ❑ **Problem:** K8s doesn't treat pods as unique, long-running instances; if a pod encounters an issue and dies, K8s replaces it for the app to minimize downtime.
- ❑ However, as pods are replaced, their internal names and IPs might change → a pod's lifetime is not reliable.
- ❑ **Solution:** An abstraction over Pods, i.e., Service, serves as the interface the app consumers interact with.
 - ❖ A service exposes a single machine name or IP address mapped to pods whose underlying names and numbers are unreliable.
 - ❖ A service ensures that, to the outside network, everything appears to be unchanged.

Service YAML Structure

❑ Common spec fields

- ❖ ports
- ❖ type

❑ spec:ports:port

- ❖ the port number on which the Service will listen for incoming traffic
- ❖ used to communicate with the Service

❑ spec:ports:protocol

- ❖ specifies the protocol used for the port
- ❖ if not specified, TCP is used by default

```
apiVersion: v1
kind: Service
metadata:
  name: http-service
spec:
  ports:
    - name: http
      port: 80
      protocol: TCP
  type: ...
```

Service type: ClusterIP

❑ Scenario: Internal Service Discovery


- ❖ client: pods inside K8s cluster
- ❖ server: pods inside K8s cluster

❑ Type-specific field: targetPort

- ❖ the port on the backend Pods to which the incoming traffic will be forwarded
- ❖ the port that your application inside the Pod is listening on

❑ Type-specific field: selector

- ❖ conditions target Pods need to satisfy
- ❖ used to match Pod labels

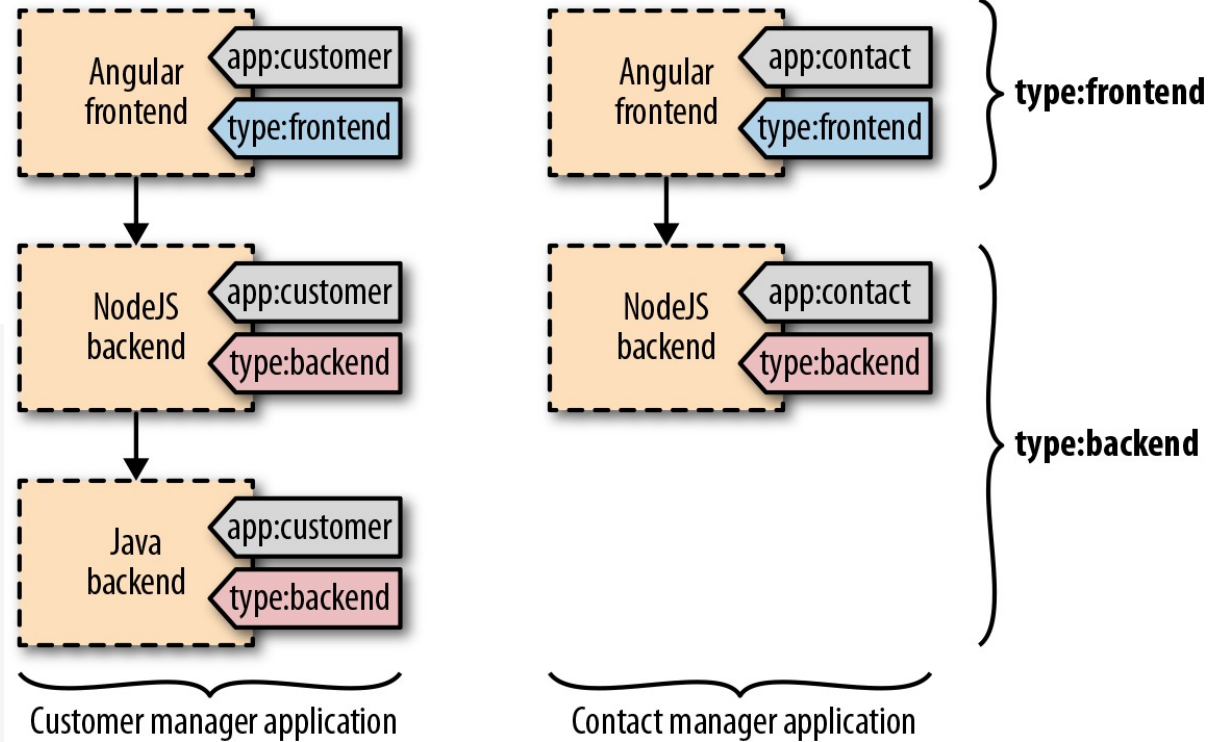


```
apiVersion: v1
kind: Service
metadata:
  name: clusterip-service
spec:
  type: ClusterIP
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: server
```

Labels



```
apiVersion: v1
kind: Pod
metadata:
  name: server-pod-v2
  labels:
    app: server
spec:
  containers:
  - name: server-container
    image: yancaobao/server-image
```



□ used as an application identity for Pods

ClusterIP - how to use?

- once a Service is created, it gets a clusterIP assigned that is accessible only from within the Kubernetes cluster, and that IP remains unchanged as long as the Service definition exists.
- when a client knows the name of the Service it wants to access, it can reach the Service by a fully qualified domain name (FQDN) such as `SERVICE_Name.default.svc.cluster.local`

```
$ kubectl port-forward service/$SERVICE_Name 8888:8080
```

Service type: ExternalName


❑ Scenario: External Service Discovery

- ❖ client: pods inside K8s cluster
- ❖ server: pods outside K8s cluster

❑ redirect connections to external IP addresses and ports by omitting the selector definition of a Service and manually creating endpoint resources

❑ Type-specific field: externalName


- ❖ CNAME record that points to the external name you specified
- ❖ try minikube using `host.minikube.internal`



```
apiVersion: v1
kind: Service
metadata:
  name: externalname-service
spec:
  type: ExternalName
  externalName: my-server.com
    - port: 80
```

Service type: NodePort

- ❑ Scenario: Service Discovery from outside K8s cluster
 - ❖ client: outside K8s cluster
 - ❖ server: pods inside K8s cluster
- ❑ Type-specific field: nodePort
 - ❖ port opened on all nodes, through which pods inside the cluster are accessed
 - ❖ has a range from 30000 to 32767



```
apiVersion: v1
kind: Service
metadata:
  name: nodeport-service
spec:
  type: NodePort
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30036
  selector:
    app: server
```