

# Kubernetes Design Patterns

SoC Summer Workshop  
Cloud Computing with Big Data

**Richard T. B. Ma**

School of Computing

National University of Singapore

# Roadmap

## ❑ Behavioral Patterns

- ❖ Stateful service -- StatefulSet
- ❖ Self-awareness -- Downward API

## ❑ Structural Patterns

- ❖ Init container
- ❖ Sidecar container

## ❑ Lifecycle Patterns

- ❖ Health probes
- ❖ Managed lifecycle

# Servers in the DevOps World

## Pets



- ❑ nonfungible servers
  - ❖ every instance is unique
- ❑ require individual care
  - ❖ repair
  - ❖ vertical scaling
- ❑ stateful, persistent and permanent

## Cattle



- ❑ identical servers
  - ❖ all instances are the same
- ❑ not need individual care
  - ❖ replaced
  - ❖ horizontal scaling
- ❑ stateless, ephemeral, and transient

# StatefulSet (STS)

- ❑ Distributed stateful apps require
  - ❖ persistent storage, identity, networking, and ordinality
  - ❖ every instance is unique & has long-lived characteristics
  - ❖ e.g., big data frameworks such as Map-Reduce
- ❑ Solution: StatefulSets provides
  - ❖ stable, unique network identifiers: each Pod in a STS gets a unique hostname based on its ordinal index.
  - ❖ stable, persistent storage: each Pod can be associated with a PersistentVolume.
  - ❖ ordered, automated rolling updates: STS manages the deployment and scaling in an ordered & deterministic fashion

# StatefulSet - how to use?

- ❑ Step 1: create a headless service
  - ❖ a ClusterIP Service without a virtual IP

- ❑ Usage case:

- ❖ direct access to the individual pods without load balancing

- ❑ How does it work?

- ❖ `headless-service.default.svc.cluster.local` will resolve to multiple IPs, one for each Pod.
  - ❖ `Pod-name.headless-service.default.svc.cluster.local` will resolve to the specific Pod's IP.

```
apiVersion: v1
kind: Service
metadata:
  name: headless-service
spec:
  clusterIP: None
  selector:
    app: server
  ports:
    - port: 80
```

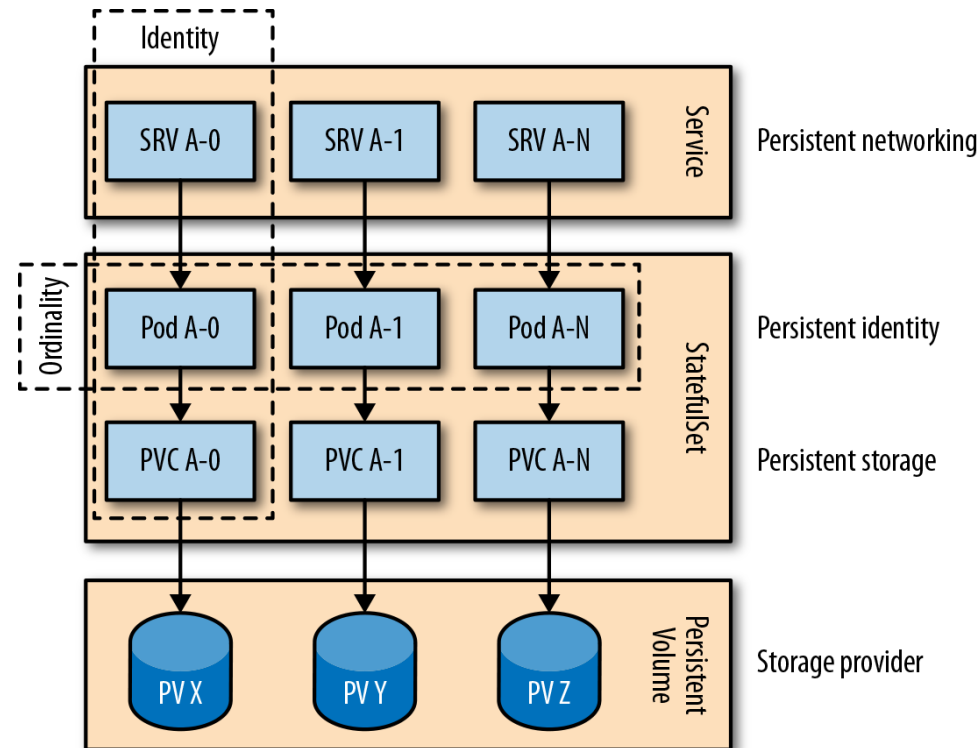
# StatefulSet - how to use?

- ❑ Step 2: create a StatefulSet
  - ❖ serviceName matches headless service
- ❑ volumeClaimTemplates mechanism
  - ❖ specifies storage requirements
  - ❖ creates PVCs on the fly during
  - ❖ allows each Pod to get its own dedicated PVC during pod creation
- ❑ In contrast, Deployment & ReplicaSet
  - ❖ use a predefined PVC, suited for using ReadOnlyMany or ReadWriteMany volumes mounted on multiple replicas
  - ❖ not suited for ReadWriteOnce volumes

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: server-statefulset
spec:
  selector:
    matchLabels:
      app: server
  serviceName: "headless-service"
  replicas: 3
  template:
    metadata:
      labels:
        app: server
    spec:
      containers:
        - name: server-container
          image: yancanmao/server-image
          ports:
            - containerPort: 80
              name: web
          volumeMounts:
            - name: www
              mountPath: /usr/share/server
  volumeClaimTemplates:
    metadata:
      name: www
    spec:
      accessModes: [ "ReadWriteOnce" ]
      resources:
        requests:
          storage: 1Gi
```

# StatefulSet - Characteristics

- ❑ STS does not manage PV
  - ❖ but manages PVCs
  - ❖ scaling up creates new Pods and associated PVCs.
  - ❖ scaling down deletes the Pods, but it does not delete any PVCs (nor PVs)
- ❑ K8s cannot free the claimed/used PV storage
  - ❖ manual deletion is needed
  - ❖ a system behavior by design



# The need of self-awareness

- ❑ Scenario: apps may need to have info about themselves and their running environment
  - ❖ runtime info: Pod's name & IP, Node's hostname
  - ❖ static info: specific resource requests & limits
  - ❖ dynamic info: annotations and labels
- ❑ Use cases:
  - ❖ log information, send metrics to a central server.
  - ❖ tune thread-pool size, GC algorithm or memory allocation based on resource limits
  - ❖ discover other pods and interact with them
- ❑ Solution: Downward API
  - ❖ allows passing metadata about the Pod to the containers and the cluster through environment variables and files



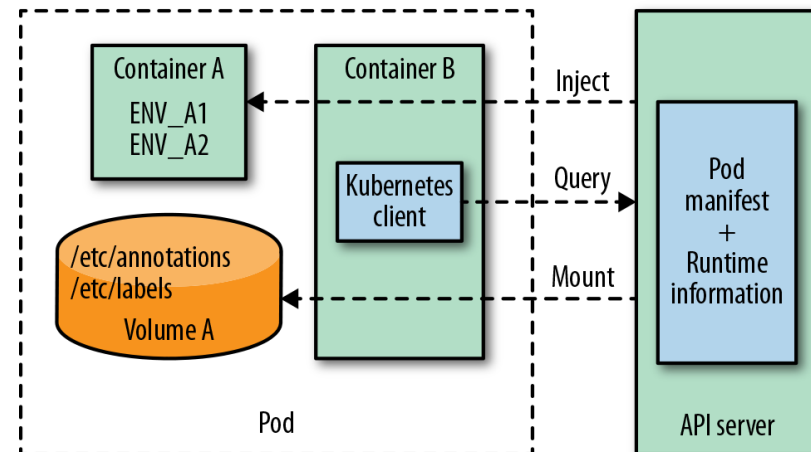
# Downward API - how does it work?

## ❑ Same mechanisms for passing data from ConfigMaps

- ❖ data is not created by developers
- ❖ specify the keys that interests us, and K8s populates the values dynamically

## ❑ Main advantage:

- ❖ metadata is injected into Pod and made available locally
- ❖ no need to use a client and interact with the API server
- ❖ nonintrusive introspection and metadata injection, remain K8s-agnostic



# Downward API - how to use?

- ❑ Import as environment variables

- ❑ `fieldPath:fieldPath` option:

- ❖ `POD_NAME`, `POD_NAMESPACE`, `POD_IP`, and `NODE_NAME` environment variables are set using the Downward API.

- ❑ `ResourceFieldRef` option:

- ❖ `CPU_LIMIT` and `MEMORY_LIMIT` are set using container resource limits.

```
apiVersion: v1
kind: Pod
metadata:
  name: downwardapi-env-pod
spec:
  containers:
  - name: nginx
    image: nginx
    env:
    - name: POD_NAME
      valueFrom:
        fieldRef:
          fieldPath: metadata.name
    - name: POD_NAMESPACE
      valueFrom:
        fieldRef:
          fieldPath: metadata.namespace
    - name: POD_IP
      valueFrom:
        fieldRef:
          fieldPath: status.podIP
    - name: NODE_NAME
      valueFrom:
        fieldRef:
          fieldPath: spec.nodeName
    - name: CPU_LIMIT
      valueFrom:
        resourceFieldRef:
          containerName: nginx
          resource: limits.cpu
          divisor: 1m
    - name: MEMORY_LIMIT
      valueFrom:
        resourceFieldRef:
          containerName: nginx
          resource: limits.memory 10
          divisor: 1Mi
```

# Downward API - how to use?

## ❑ Import as a volume

- ❖ downwardAPI type of volume
- ❖ all information written into files
- ❖ all the labels and annotations retrieved as files, not for EnvVar

## ❑ Available information:

<https://kubernetes.io/docs/concepts/workloads/pods/downward-api/>

## ❑ Limitations of downward API:

- ❖ limited info, some can only be accessed by one method

```
apiVersion: v1
kind: Pod
metadata:
  name: downwardapi-volume-pod
spec:
  containers:
  - name: nginx
    image: nginx
    volumeMounts:
    - name: downward-api-volume
      mountPath: /etc/downward
  volumes:
  - name: downward-api-volume
    downwardAPI:
      items:
      - path: labels
        fieldRef:
          fieldPath: metadata.labels
      - path: annotations
        fieldRef:
          fieldPath: metadata.annotations
      - path: cpu_limit
        resourceFieldRef:
          containerName: nginx
          resource: limits.cpu
          divisor: 1m
      - path: memory_limit
        resourceFieldRef:
          containerName: nginx
          resource: limits.memory
          divisor: 1Mi
```