

Kubernetes Design Patterns

SoC Summer Workshop
Cloud Computing with Big Data

Richard T. B. Ma

School of Computing

National University of Singapore

Roadmap

❑ Behavioral Patterns

- ❖ Stateful service -- StatefulSet
- ❖ Self-awareness -- Downward API

❑ Structural Patterns

- ❖ Init containers
- ❖ Sidecar containers

❑ Lifecycle Patterns

- ❖ Health probes
- ❖ Managed lifecycle

Servers in the DevOps World

Pets



- ❑ nonfungible servers
 - ❖ every instance is unique
- ❑ require individual care
 - ❖ repair
 - ❖ vertical scaling
- ❑ stateful, persistent and permanent

Cattle



- ❑ identical servers
 - ❖ all instances are the same
- ❑ not need individual care
 - ❖ replaced
 - ❖ horizontal scaling
- ❑ stateless, ephemeral, and transient

StatefulSet (STS)

- ❑ Distributed stateful apps require
 - ❖ persistent storage, identity, networking, and ordinality
 - ❖ every instance is unique & has long-lived characteristics
 - ❖ e.g., big data frameworks such as Map-Reduce
- ❑ Solution: StatefulSets provides
 - ❖ stable, unique network identifiers: each Pod in a STS gets a unique hostname based on its ordinal index.
 - ❖ stable, persistent storage: each Pod can be associated with a PersistentVolume.
 - ❖ ordered, automated rolling updates: STS manages the deployment and scaling in an ordered & deterministic fashion

StatefulSet - how to use?

- ❑ Step 1: create a headless service
 - ❖ a ClusterIP Service without a virtual IP

- ❑ Usage case:

- ❖ direct access to the individual pods without load balancing

- ❑ How does it work?

- ❖ `headless-service.default.svc.cluster.local` will resolve to multiple IPs, one for each Pod.
 - ❖ `Pod-name.headless-service.default.svc.cluster.local` will resolve to the specific Pod's IP.

```
apiVersion: v1
kind: Service
metadata:
  name: headless-service
spec:
  clusterIP: None
  selector:
    app: server
  ports:
    - port: 80
```

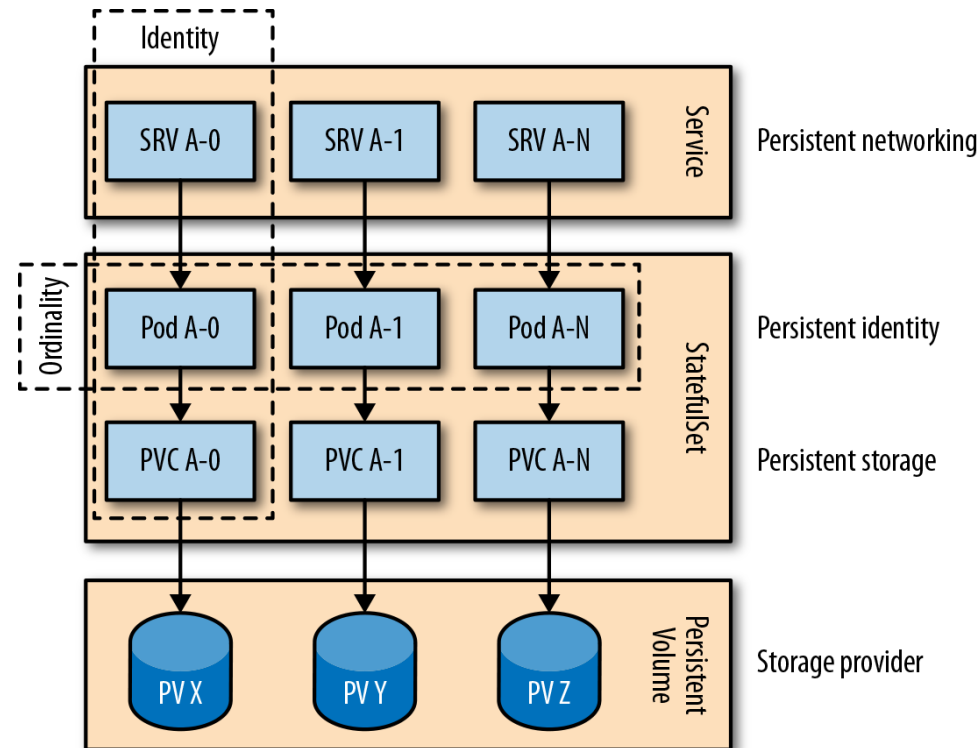
StatefulSet - how to use?

- ❑ Step 2: create a StatefulSet
 - ❖ serviceName matches headless service
- ❑ volumeClaimTemplates mechanism
 - ❖ specifies storage requirements
 - ❖ creates PVCs on the fly during
 - ❖ allows each Pod to get its own dedicated PVC during pod creation
- ❑ In contrast, Deployment & ReplicaSet
 - ❖ use a predefined PVC, suited for using ReadOnlyMany or ReadWriteMany volumes mounted on multiple replicas
 - ❖ not suited for ReadWriteOnce volumes

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: server-statefulset
spec:
  selector:
    matchLabels:
      app: server
  serviceName: "headless-service"
  replicas: 3
  template:
    metadata:
      labels:
        app: server
    spec:
      containers:
        - name: server-container
          image: yancanmao/server-image
          ports:
            - containerPort: 80
              name: web
          volumeMounts:
            - name: www
              mountPath: /usr/share/server
  volumeClaimTemplates:
    metadata:
      name: www
    spec:
      accessModes: [ "ReadWriteOnce" ]
      resources:
        requests:
          storage: 1Gi
```

StatefulSet - Characteristics

- ❑ STS does not manage PV
 - ❖ but manages PVCs
 - ❖ scaling up creates new Pods and associated PVCs.
 - ❖ scaling down deletes the Pods, but it does not delete any PVCs (nor PVs)
- ❑ K8s cannot free the claimed/used PV storage
 - ❖ manual deletion is needed
 - ❖ a system behavior by design



The need of self-awareness

- ❑ Scenario: apps may need to have info about themselves and their running environment
 - ❖ runtime info: Pod's name & IP, Node's hostname
 - ❖ static info: specific resource requests & limits
 - ❖ dynamic info: annotations and labels
- ❑ Use cases:
 - ❖ log information, send metrics to a central server.
 - ❖ tune thread-pool size, GC algorithm or memory allocation based on resource limits
 - ❖ discover other pods and interact with them
- ❑ Solution: Downward API
 - ❖ allows passing metadata about the Pod to the containers and the cluster through environment variables and files

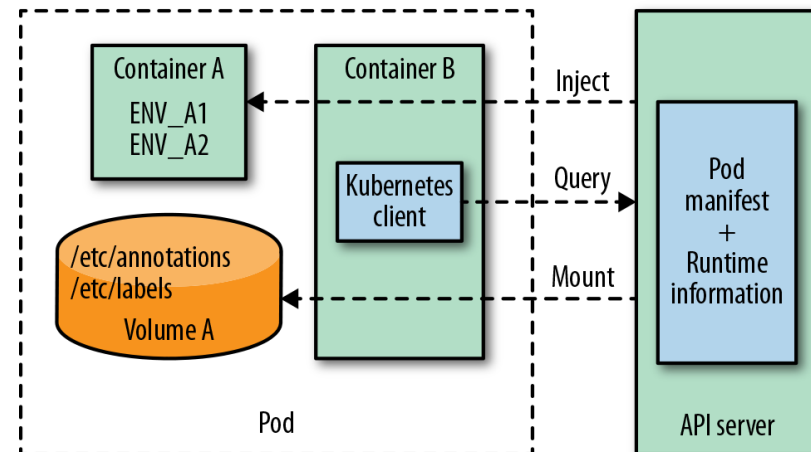
Downward API - how does it work?

❑ Same mechanisms for passing data from ConfigMaps

- ❖ data is not created by developers
- ❖ specify the keys that interests us, and K8s populates the values dynamically

❑ Main advantage:

- ❖ metadata is injected into Pod and made available locally
- ❖ no need to use a client and interact with the API server
- ❖ nonintrusive introspection and metadata injection, remain K8s-agnostic



Downward API - how to use?

- ❑ Import as environment variables

- ❑ `fieldPath:fieldPath` option:

- ❖ `POD_NAME`, `POD_NAMESPACE`, `POD_IP`, and `NODE_NAME` environment variables are set using the Downward API.

- ❑ `ResourceFieldRef` option:

- ❖ `CPU_LIMIT` and `MEMORY_LIMIT` are set using container resource limits.

```
apiVersion: v1
kind: Pod
metadata:
  name: downwardapi-env-pod
spec:
  containers:
  - name: nginx
    image: nginx
    env:
    - name: POD_NAME
      valueFrom:
        fieldRef:
          fieldPath: metadata.name
    - name: POD_NAMESPACE
      valueFrom:
        fieldRef:
          fieldPath: metadata.namespace
    - name: POD_IP
      valueFrom:
        fieldRef:
          fieldPath: status.podIP
    - name: NODE_NAME
      valueFrom:
        fieldRef:
          fieldPath: spec.nodeName
    - name: CPU_LIMIT
      valueFrom:
        resourceFieldRef:
          containerName: nginx
          resource: limits.cpu
          divisor: 1m
    - name: MEMORY_LIMIT
      valueFrom:
        resourceFieldRef:
          containerName: nginx
          resource: limits.memory
          divisor: 1Mi
```

Downward API - how to use?

❑ Import as a volume

- ❖ downwardAPI type of volume
- ❖ all information written into files
- ❖ all the labels and annotations retrieved as files, not for EnvVar

❑ Available information:

<https://kubernetes.io/docs/concepts/workloads/pods/downward-api/>

❑ Limitations of downward API:

- ❖ limited info, some can only be accessed by one method

```
apiVersion: v1
kind: Pod
metadata:
  name: downwardapi-volume-pod
spec:
  containers:
  - name: nginx
    image: nginx
    volumeMounts:
    - name: downward-api-volume
      mountPath: /etc/downward
  volumes:
  - name: downward-api-volume
    downwardAPI:
      items:
      - path: labels
        fieldRef:
          fieldPath: metadata.labels
      - path: annotations
        fieldRef:
          fieldPath: metadata.annotations
      - path: cpu_limit
        resourceFieldRef:
          containerName: nginx
          resource: limits.cpu
          divisor: 1m
      - path: memory_limit
        resourceFieldRef:
          containerName: nginx
          resource: limits.memory
          divisor: 1Mi
```

Roadmap

□ Behavioral Patterns

- ❖ Stateful service -- StatefulSet
- ❖ Self-awareness -- Downward API


□ Structural Patterns

- ❖ Init containers
- ❖ Sidecar containers

□ Lifecycle Patterns

- ❖ Health probes
- ❖ Managed lifecycle

Init Containers



- ❑ Initialization is common in software development
 - ❖ e.g., constructors in OOP: executed only once at the beginning of the creation of a class instance
 - ❖ how could we appropriately initialize Pods?
- ❑ Solution: use **init containers** in the Pod definition
 - ❖ two types of containers: init and application containers
 - ❖ init containers are executed in a sequence, one by one
 - ❖ all of them have to terminate successfully before the application containers are started up
 - ❖ if an init container fails, the whole Pod is restarted, causing all init containers to run again; thus, to prevent side effects, init containers need to be idempotent

Init Containers - how to use?

- ❑ Enable separation of concerns
 - ❖ app engineers focus on app logic
 - ❖ develop engineers focus on configuration and initialization
- ❑ Init containers are typically small, run quickly, and complete successfully
 - ❖ except when used to delay the start of a Pod while waiting for a dependency
- ❑ Init containers have separate environments, but can share volumes with app containers

```
apiVersion: v1
kind: Pod
metadata:
  name: www
  labels:
    app: www
spec:
  initContainers:
    - name: download
      image: bitnami/git
      # Clone an HTML page to be served
      command:
        - git
        - clone
        - https://github.com/mdn/beginner-html-site-scripted
        - /var/lib/data
      # Shared volume with main container
      volumeMounts:
        - mountPath: /var/lib/data
          name: source
  containers:
    # Simple static HTTP server for serving these pages
    - name: run
      image: docker.io/centos/httpd
      ports:
        - containerPort: 80
      # Shared volume with main container
      volumeMounts:
        - mountPath: /var/www/html
          name: source
  volumes:
    - emptyDir: {}
      name: source
```

Sidecar containers -- motivations




- ❑ A Pod can run multiple containers. How should we organize the application logics into containers?
 - ❖ think about a natural boundary for a unit of functionality with a distinct runtime, release cycle, API, develop team.
- ❑ A proper container behaves like a single process
 - ❖ solves one problem and does it well
 - ❖ is created with the idea of replaceability and reuse
- ❑ Single-purpose reusable containers require
 - ❖ ways of extending the functionality of a container
 - ❖ a means for collaboration among containers.

Sidecar container example

□ an HTTP server and a Git synchronizer

- ❖ the HTTP server focuses only on serving files over HTTP and does not know how and where the files are coming from.
- ❖ the Git synchronizer's only goal is to sync data from a Git server to the local filesystem and does not care what happens to the files once synced.



```
apiVersion: v1
kind: Pod
metadata:
  name: web-app
spec:
  containers:
    # Main container is a stock httpd serving from /var/www/html
    - name: app
      image: centos/httpd
      ports:
        - containerPort: 80
      volumeMounts:
        - mountPath: /var/www/html
          name: source
    # Sidecar poll every minute a given repository with git
    - name: poll
      image: bitnami/git
      env:
        - name: SOURCE_REPO
          value: https://github.com/mdn/beginner-html-site-scripted
      command: [ "sh", "-c" ]
      args:
        - |
          git clone $(SOURCE_REPO) .
          while true; do
            sleep 60
            git pull
          done
      workingDir: /var/lib/data
      volumeMounts:
        - mountPath: /var/lib/data
          name: source
  volumes:
    # The shared directory for holding the files
    - emptyDir: {}
      name: source
```



Sidecar containers - use cases

1. **Data Synchronization:** Syncing data between the main container and an external storage or database.
2. **Logging & Monitoring:** Collecting logs and metrics from the main container and sending them to external logging or monitoring systems.
3. **Service Mesh Proxies:** Injecting sidecar proxies (e.g., Envoy in Istio) to manage network traffic and enhance security, observability, and reliability.
4. **Configuration & Secrets Management:** Updating configs or secrets dynamically without restarting the main container.
5. **Security:** Providing security features like authentication, encryption, or token refresh.

Sidecars patterns: Adapter and Ambassador

- ❑ Scenario 1: your application is complicated: heterogeneous components in system from multiple teams using different technologies
 - ❖ how to utilize external service for all? it boils down to:
 - ❖ how to make it conform to a consistent, unified interface with a standardized and normalized format that can be consumed by the outside world?
- ❑ Scenario 2: external services are complicated: heterogeneous external alternative choices, e.g., memory cache and database storage
 - ❖ how to hide complexity and providing a unified interface for accessing services outside the Pod?


Adapter pattern



❑ example: when leveraging external monitoring tools

- ❖ distributed software components written in different languages may not have the same capabilities
- ❖ they also may not expose metrics in the same format expected by the monitoring tool
- ❖ Solution: every service represented by a Pod, in addition to the main application container, would have another container that knows how to read the custom application-specific metrics and expose them in a generic format understandable by the monitoring tool.

Ambassador pattern



❑ example: when leveraging heterogeneous storage

- ❖ an application needs to store data using different types of external storage based on different conditions
- ❖ e.g., local memory and remote data store
- ❖ Solution: a sidecar container acts as a proxy and decouples the main Pod from directly accessing external dependencies.

Roadmap

□ Behavioral Patterns

- ❖ Stateful service -- StatefulSet
- ❖ Self-awareness -- Downward API

□ Structural Patterns

- ❖ Init containers
- ❖ Sidecar containers


□ Lifecycle Patterns

- ❖ Health probes
- ❖ Managed lifecycle

Health probes

□ How can an application communicate its health state to Kubernetes?

- ❖ some container observability options are as follows:



□ kubelet constantly performs a **process health check** on the container processes

- ❖ container is restarted if its processes are not running
- ❖ however, not sufficient to decide about the health of an app; e.g., a Java app may throw an `OutOfMemoryError` and still have the JVM process running.

Liveness probes

- ❑ Checks performed by kubelet
 - ❖ ask container to confirm its healthiness
 - ❖ have the checks probed from outside
- ❑ Flexible methods:
 - ❖ HTTP GET request to the container IP address and expects a successful HTTP response code between 200 and 399.
 - ❖ A TCP Socket probe assumes a successful TCP connection.
 - ❖ an Exec probe executes an arbitrary command in the container kernel namespace and expects a successful exit code (0).

```
apiVersion: v1
kind: Pod
metadata:
  name: livenessprobe-pod
spec:
  containers:
    - image: yancanmao/server-image
      name: server-container
      livenessProbe:
        httpGet:
          path: /healthy
          port: 8080
          initialDelaySeconds: 5
          timeoutSeconds: 1
          periodSeconds: 10
          failureThreshold: 3
        ports:
          - containerPort: 8080
            name: http
            protocol: TCP
```


Readiness probes

```
readinessProbe:  
  exec:  
    command:  
      - cat  
      - /tmp/healthy  
  initialDelaySeconds: 5  
  periodSeconds: 5
```

- ❑ Scenario: although a container is not healthy, restarting it may not help either, for example
 - ❖ when a container is still starting up and not ready to handle any requests yet
 - ❖ or a container is overloaded, and its latency is increasing, better shield it from additional load for a while.
- ❑ Different corrective action from liveness probe
 - ❖ a failed probe causes the container to be removed from the service endpoint and not receive any new traffic.
 - ❖ process health and liveness checks are intended to recover from failure by restarting; readiness check buys time for your app and expects it to recover by itself

Managed Lifecycle

- ❑ Containerized apps managed by cloud-native platforms have no control over their lifecycle
 - ❖ listen to the events emitted by the managing platform and adapt their lifecycles accordingly.
 - ❖ however, health-check APIs are read-only endpoints the platform uses to extract info from the app.
- ❑ Container can listen to and react to the following events emitted by the platform if desired:



SIGTERM and SIGKILL signals

- ❑ When K8s decides to shut down a container
 - ❖ the container receives a SIGTERM signal
 - ❖ the app should shut down as quickly as possible: some apps may have to complete in-flight requests, release open connections, and clean up temp files, which can take a longer time.
- ❑ After a grace period, if not shut down yet
 - ❖ the container receives a SIGKILL signal
 - ❖ configurable grace period at pod-level
`.spec.terminationGracePeriodSeconds`
 - ❖ the container will be shut down forcefully
- ❑ It is up to the app to respond to signals in its logic.

```
spec:
  containers:
  - image: yancanmao/server-image
    name: server-container
    lifecycle:
      postStart:
        exec:
          command: ["/bin/sh", "-c", "echo Hello from the postStart hook > /var/log/poststart.log"]
      preStop:
        exec:
          command: ["/bin/sh", "-c", "echo Goodbye from the preStop hook > /var/log/prestop.log"]
```

postStart & preStop hooks

- ❑ Lifecycle hooks provided by K8s,
 - ❖ invocation mechanisms: httpGet and exec
 - ❖ both postStart and preStop are blocking calls
- ❑ The postStart command is executed after a container is created, asynchronously with the container's process
 - ❖ the container status remains Waiting until the postStart handler completes, which in turn keeps the Pod status in the Pending state
 - ❖ can be used to delay the startup state of the container while giving time to the main container process to initialize, or prevent a container from starting when the Pod does not fulfill certain preconditions
- ❑ The preStop action must complete before the call to delete the container is sent to the container runtime,
 - ❖ a convenient alternative to a SIGTERM signal