

# Kubernetes Resources & Extensions

SoC Summer Workshop  
Cloud Computing with Big Data

**Richard T. B. Ma**

School of Computing

National University of Singapore

# Roadmap

- Kubernetes API Resources

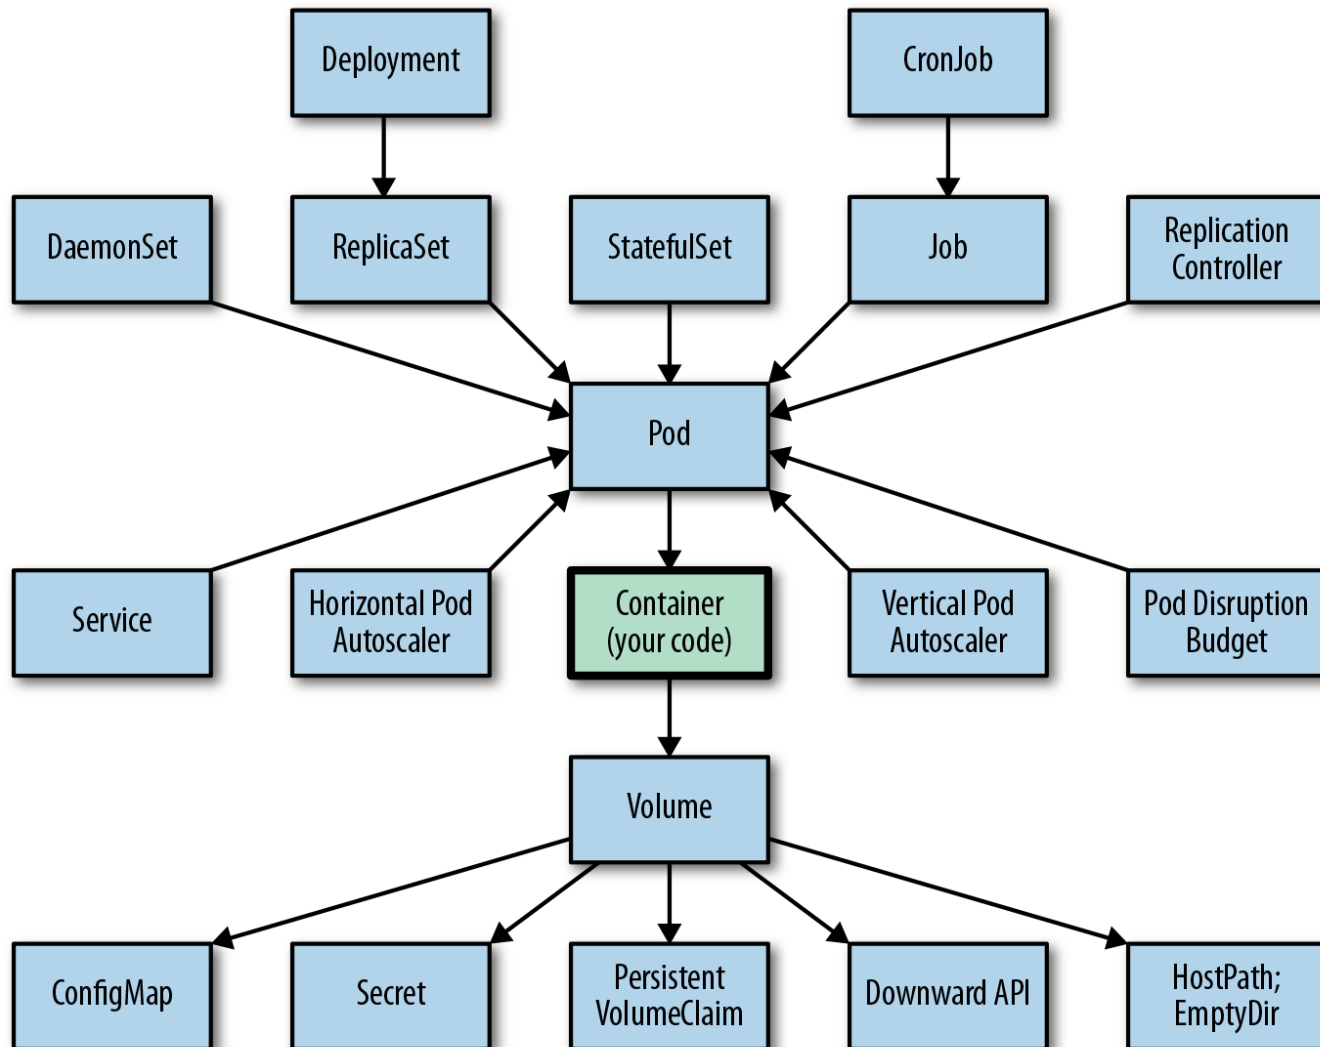
- Kubernetes Extensions

  - ❖ Controllers

  - ❖ Custom Resources

  - ❖ The Operator Pattern

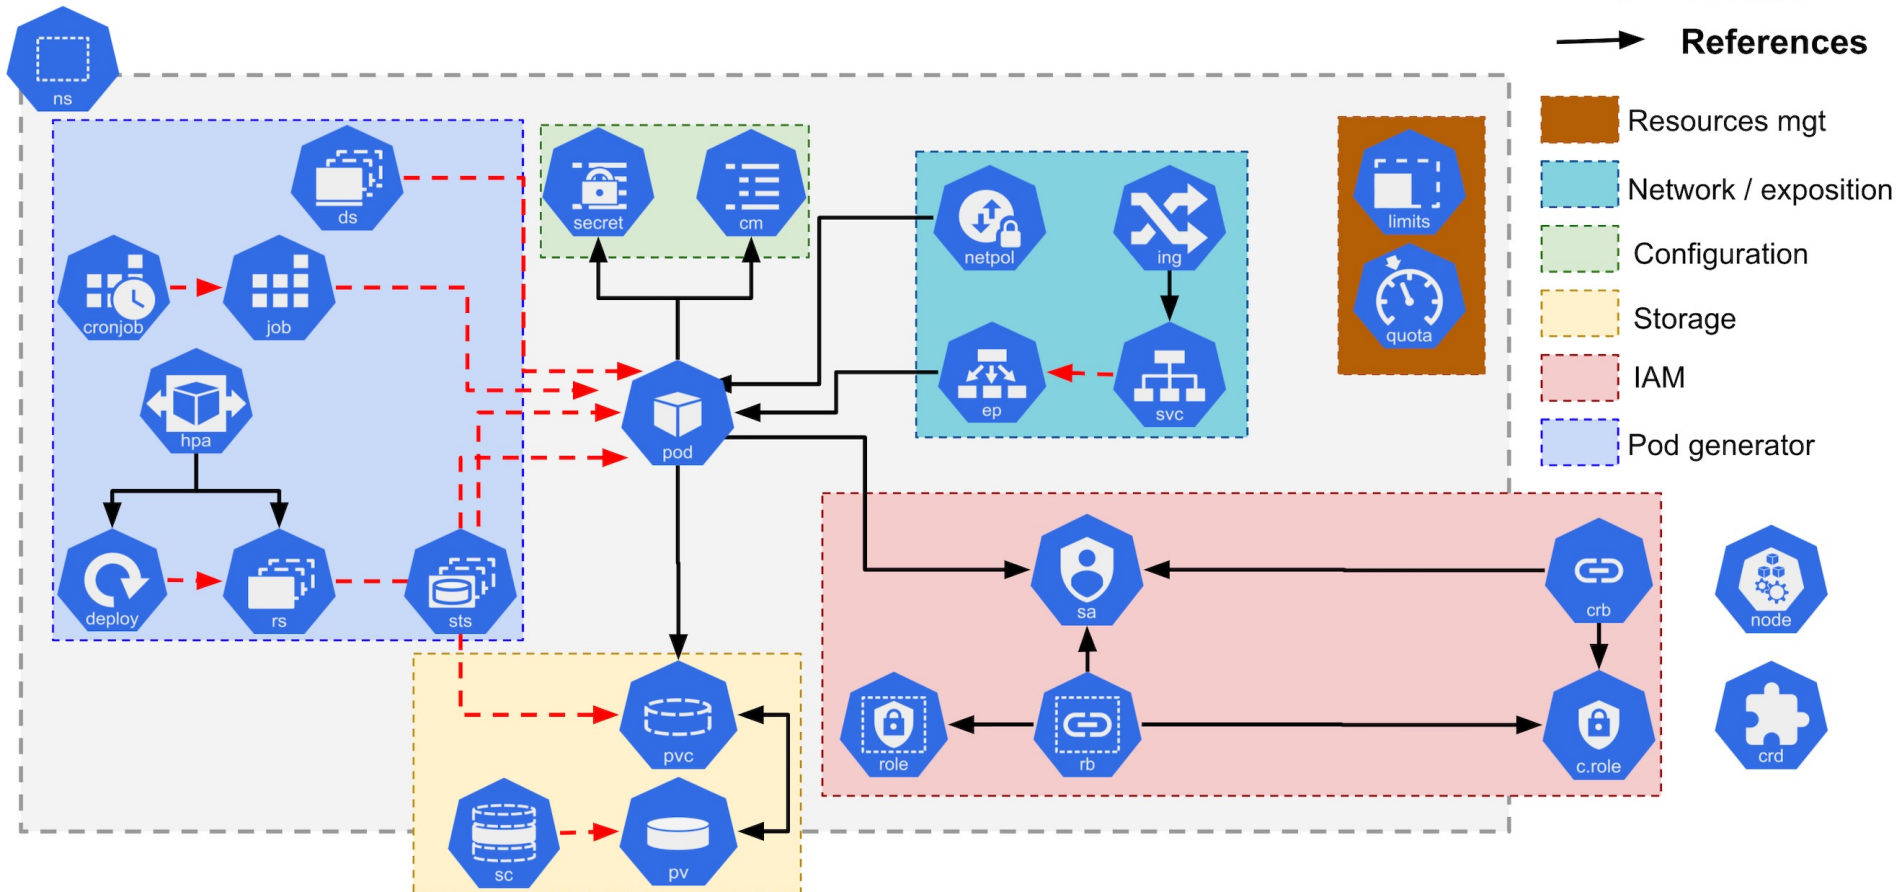
# A view of K8s for app developers



# Kubernetes Resources Map

---> Creates  
--> References

Resources mgt  
Network / exposition  
Configuration  
Storage  
IAM  
Pod generator

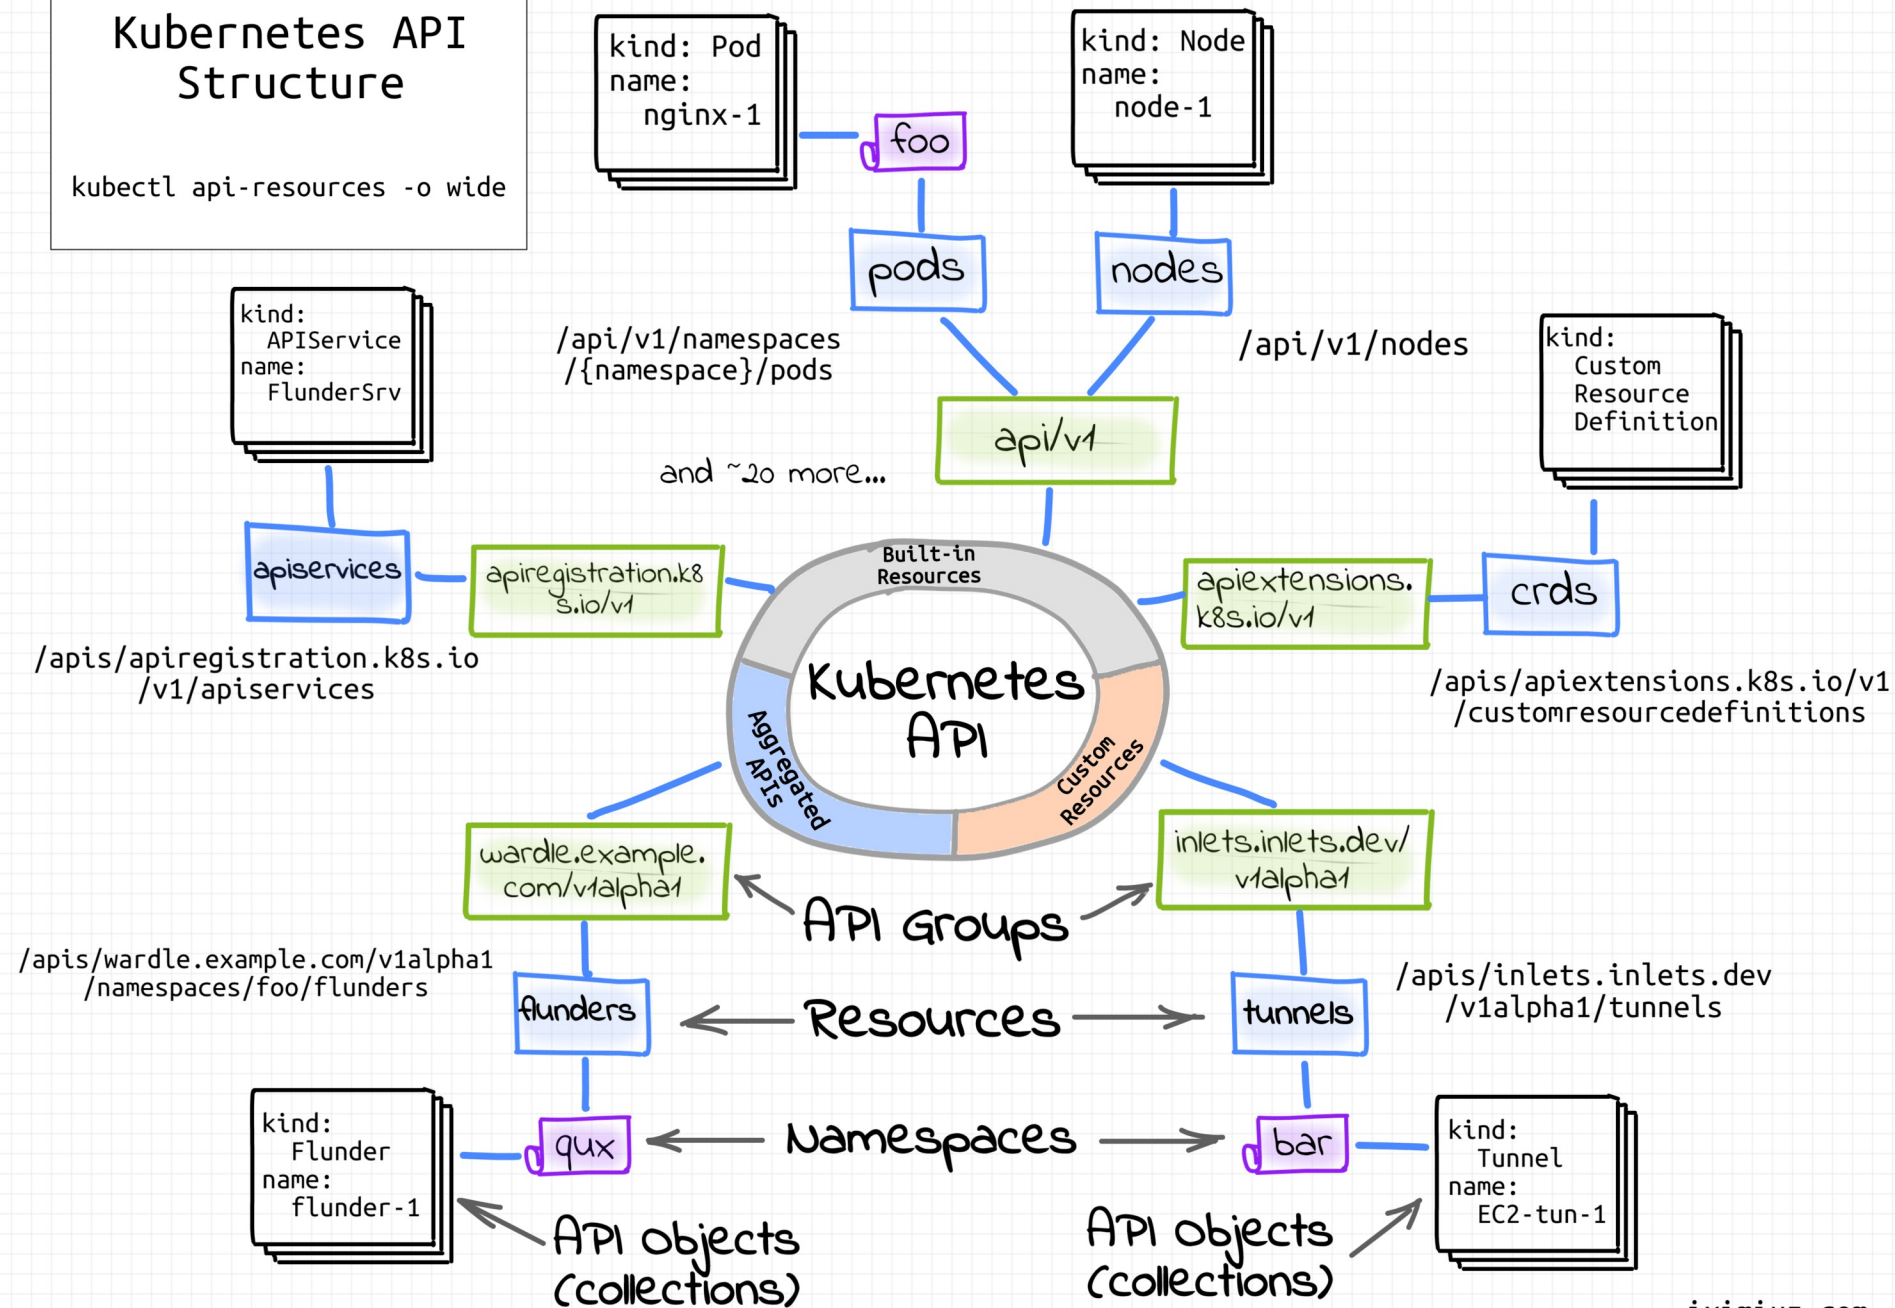


# Kubernetes Resources

- ❑ Resources refer to the API endpoints that allow you to interact with the objects in the cluster.
- ❑ Resources vs Objects
  - ❖ Objects are persistent entities in K8s
    - represent an *intent* (desired state)
    - and the *status* (actual state) of the cluster
  - ❖ Resource is more general concept
    - computing components and entities that can be managed within a Kubernetes cluster
    - include both object and non-object resources
    - RESTful like APIs via the API Server

# Kubernetes API Structure

kubectl api-resources -o wide



# Types of API Resources

## ❑ Built-in Resources

- ❖ include all native objects
- ❖ has a special CustomResourceDefinition kind

## ❑ Custom Resources

- ❖ created by CRDs
- ❖ the Operator pattern

## ❑ Aggregated APIs

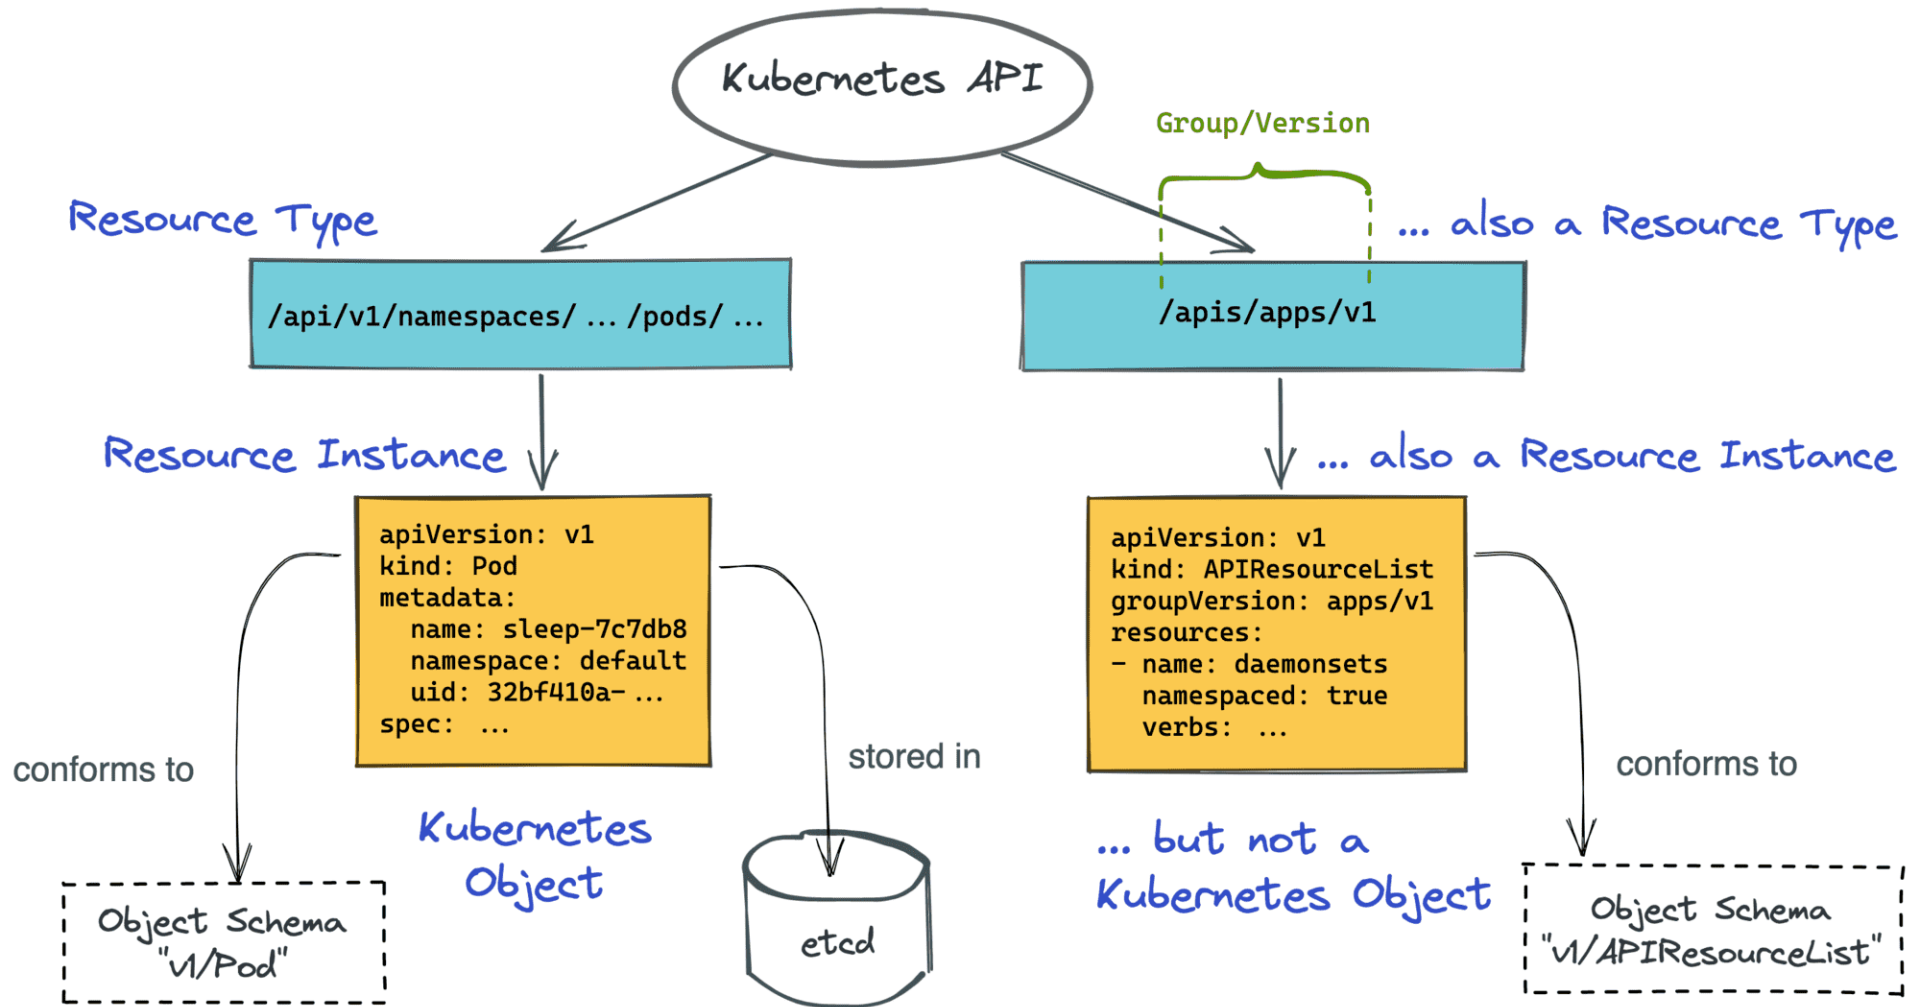
- ❖ extend the functionality of API server
- ❖ create new API endpoints via APIService kind
- ❖ hosted outside of the main K8s API server

# API Resource Groups

- ❑ Resources are bound together in API groups
  - ❖ each group may have one or more versions that evolve independent of other API groups,
  - ❖ each version within the group has one or more resources.
- ❑ Group names are typically in domain name form
  - ❖ the Kubernetes project reserves use of the empty group, all single word names ("extensions", "apps"), and any group name ending in "\*.k8s.io" for its sole use.
- ❑ When choosing a group name, recommend a subdomain your group or organization owns
  - ❖ e.g., "widget.mycompany.com".



# Object vs Non-Object Resources



# Non-Object Resources

- ❑ Resources not treated as Object:
  - ❖ **Nodes:** the machines (physical or virtual) that form the cluster.
  - ❖ **Namespaces:** provide a way to partition a cluster into multiple virtual clusters and used to organize and isolate resources within a cluster.
  - ❖ **ConfigMaps and Secrets:** are used to store configuration data and sensitive information, respectively, in a Kubernetes cluster.
- ❑ Resources managed by the Kubernetes control plane.
  - ❖ **Persistent Volumes (PVs) and Persistent Volume Claims (PVCs):** are used to manage persistent storage. PVs represent storage volumes provisioned by the admin; PVCs are requests for storage made by Pods.
  - ❖ **Service Accounts:** provide an identity for Pods running in a Kubernetes cluster. They are used by Pods to authenticate and authorize with other cluster components, such as the API server.
  - ❖ **Cluster Roles and Cluster Role Bindings:** are used to define sets of permissions (RBAC) for accessing cluster-wide resources. They are similar to roles and role bindings but operate at the cluster level.

# RESTful style API

- ❑ Create, delete, retrieve, or update a description of an object via the standard HTTP verbs (POST, PUT, DELETE, GET)
  - ❖ APIs preferentially accept and return JSON.
  - ❖ has a schema, identified by kind and apiVersion fields.
  - ❖ also exposes additional endpoints for non-standard verbs.
- ❑ Create a local proxy that acts as an intermediary between local machine and the API server:

```
$ kubectl proxy --port=8001
```

  - ❖ allows access the API securely from local machine without complex authentication settings. Now try:

```
$ curl -X GET http://localhost:8001/api/v1/namespaces/<namespace>/pods/<pod-name>
```

# Roadmap

- Kubernetes API Resources

- Kubernetes Extensions

- ❖ Controllers
- ❖ Custom Resources
- ❖ The Operator Pattern

# Controller as a design pattern

- ❑ Conceptually, it is easily a loop of the following
  - ❖ obtain runtime status
  - ❖ obtain current spec
  - ❖ check difference, make changes for both to converge
- ❑ Can be considered as a design pattern
  - ❖ an example on ConfigMap:  
<https://github.com/k8spatterns/examples/tree/main/advanced/Controller>
- ❑ Simple controllers can be built on native objects
  - ❖ as an extension to enable new logics and functionalities
  - ❖ typically run as a Deployment with just one replica
  - ❖ but invisible to the users of the cluster

# Custom Resource Definition (CRD)

❑ used to create custom resources, i.e., extensions of the K8s API.

- ❖ `spec.scope`: Namespaced or Cluster
- ❖ `spec.group`: the API group
- ❖ `spec.versions.served`: by API server?
- ❖ `spec.versions.storage`: by etcd cluster?
- ❖ `spec.versions.schema`: open standard

❑ create a custom resource (CR):

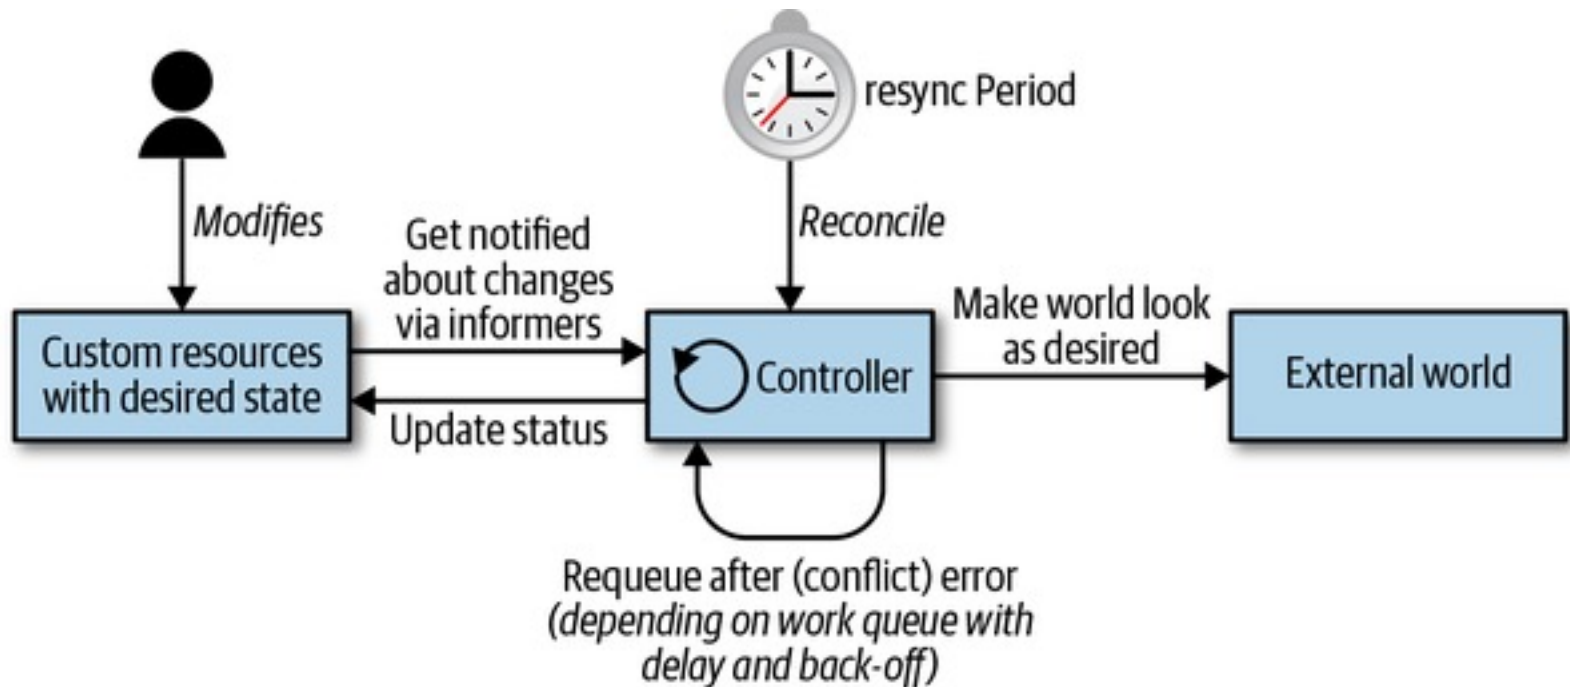
❑ controller?

```
apiVersion: example.com/v1
kind: Widget
metadata:
  name: my-widget
spec:
  size: large
  color: red
```

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: widgets.example.com
spec:
  group: example.com
  versions:
    - name: v1
      served: true
      storage: true
      schema:
        openAPIV3Schema:
          type: object
          properties:
            spec:
              type: object
              properties:
                size:
                  type: string
                color:
                  type: string
  scope: Namespaced
  names:
    plural: widgets
    singular: widget
    kind: Widget
    shortNames:
      - wgt
```

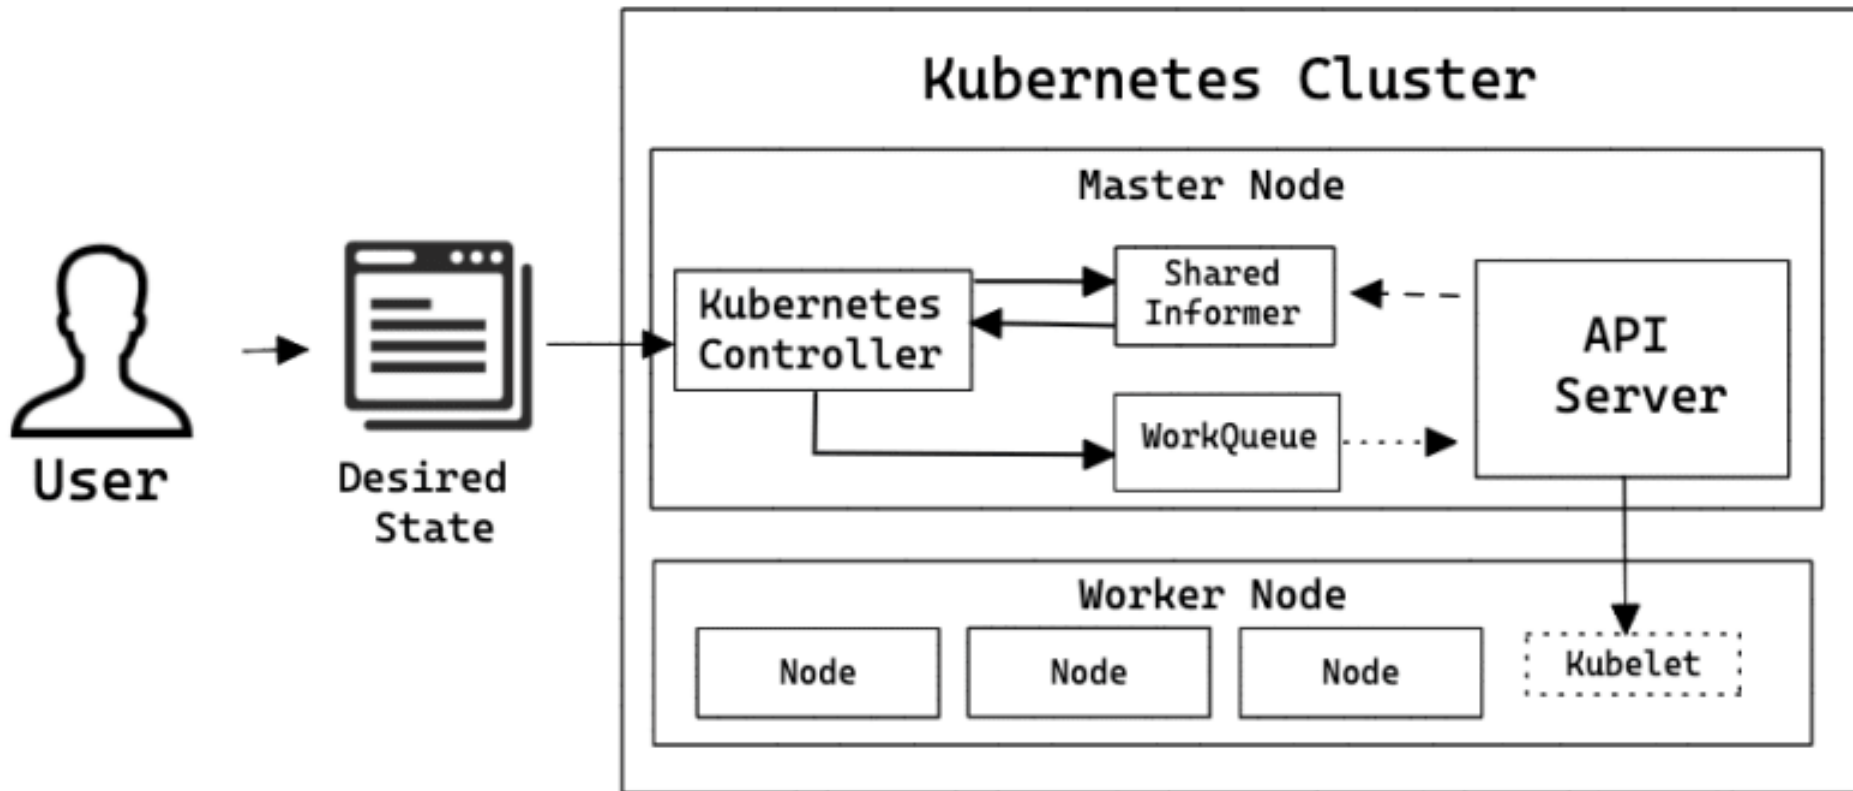
# Kubernetes Controllers

- ❑ The detailed architecture of Kubernetes native controllers is more complicated.



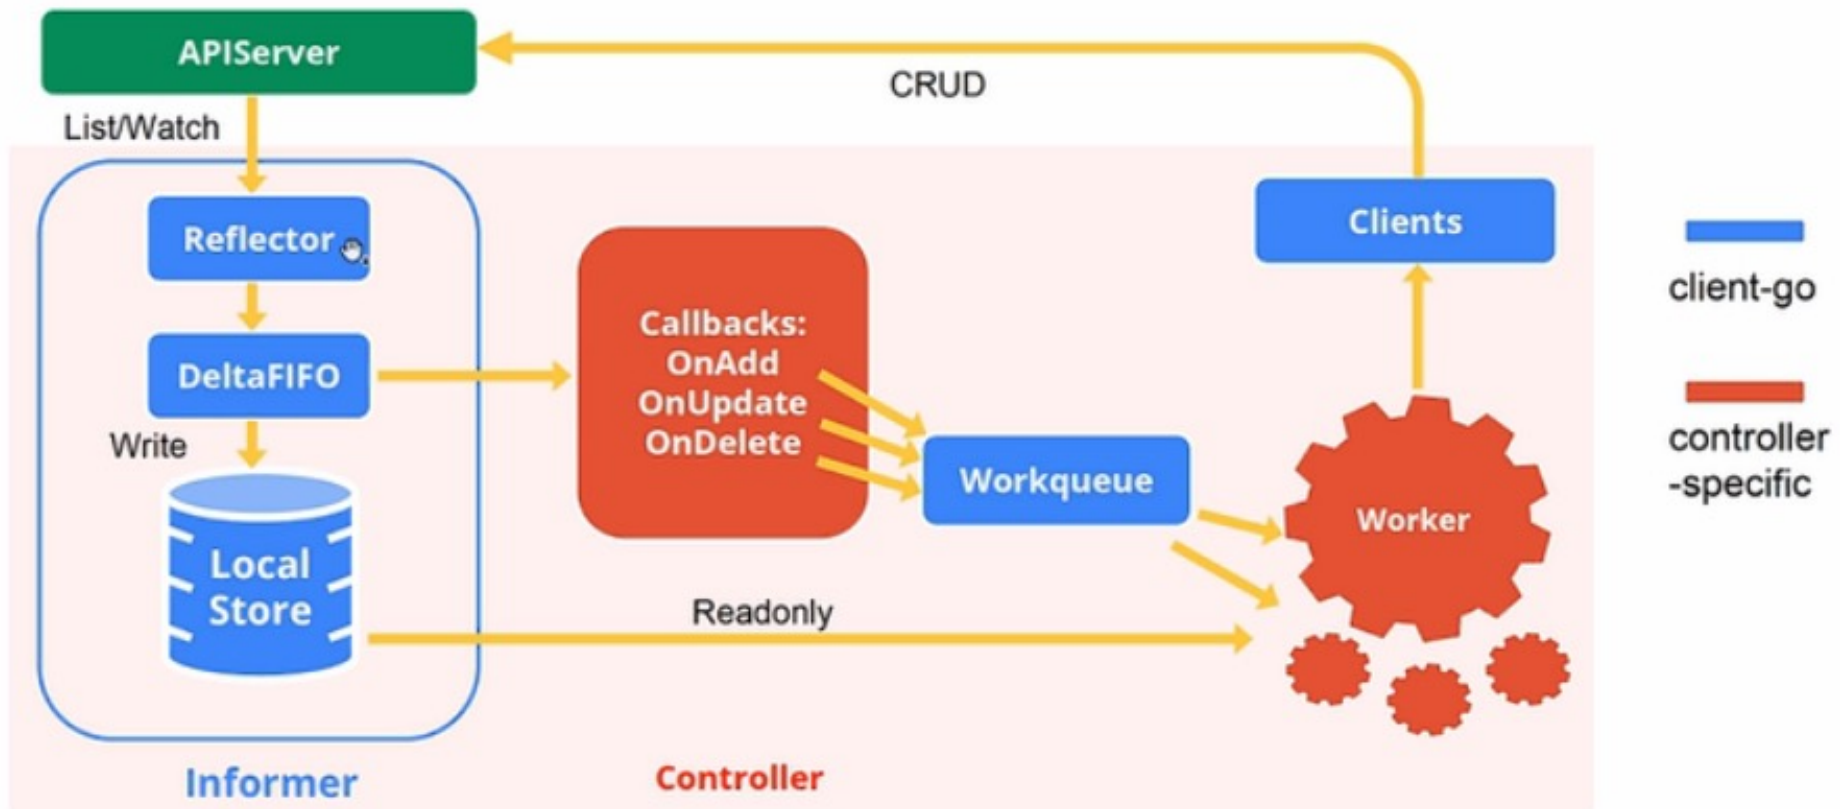
# Kubernetes Controllers

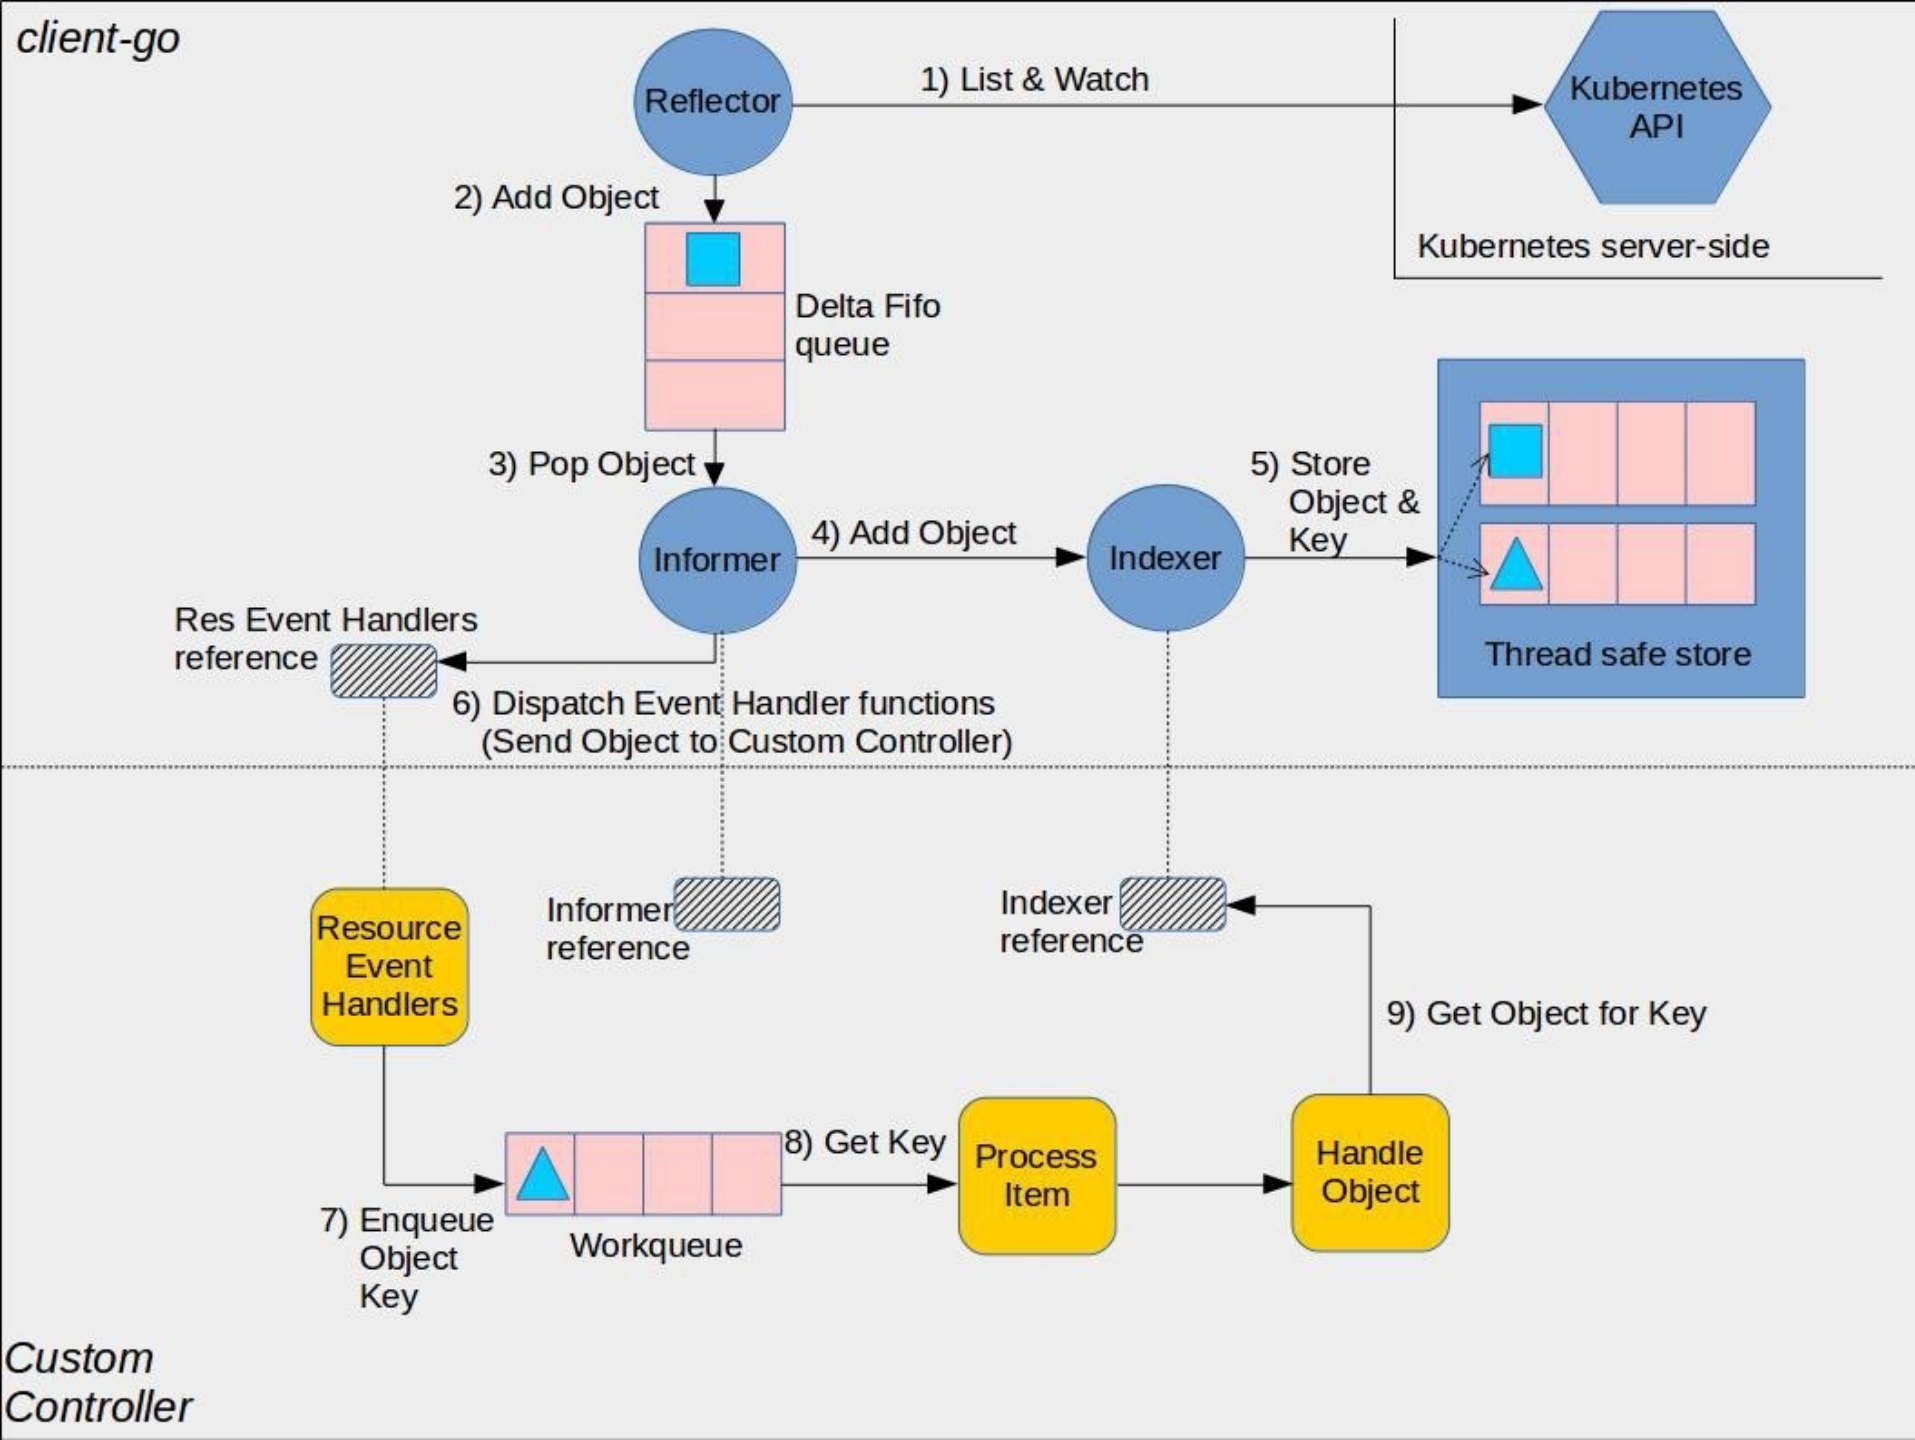
## Working of a Kubernetes Controller





# General pattern of a Kubernetes controller





# A custom controller written in Go

- ❑ watches for changes to ConfigMaps and logs these changes and leverages the client-go library
- ❑ step 1: init a new *GO* module:

```
$ mkdir configmap-controller  
$ cd configmap-controller  
$ go mod init configmap-controller  
$ go get k8s.io/client-go@kubernetes-1.26.1
```
- ❑ step 2: write the controller code `main.go`
- ❑ step 3: build and run controller:

```
$ go build -o configmap-controller .  
$ ./configmap-controller
```

# main.go

## □ important libs:

- ❖ apimachinery
- ❖ client-go

```
package main

import (
    "context"
    "fmt"
    "time"

    "k8s.io/apimachinery/pkg/fields"
    "k8s.io/apimachinery/pkg/util/runtime"
    "k8s.io/apimachinery/pkg/util/wait"
    "k8s.io/client-go/informers"
    "k8s.io/client-go/kubernetes"
    "k8s.io/client-go/tools/cache"
    "k8s.io/client-go/tools/clientcmd"
    "k8s.io/client-go/util/workqueue"
    "k8s.io/klog/v2"
)

func main() {
    // Set up Kubernetes client
    kubeconfig := clientcmd.RecommendedHomeFile
    config, err := clientcmd.BuildConfigFromFlags("", kubeconfig)
    if err != nil {
        klog.Fatalf("Error building kubeconfig: %s", err.Error())
    }

    clientset, err := kubernetes.NewForConfig(config)
    if err != nil {
        klog.Fatalf("Error building Kubernetes clientset: %s", err.Error())
    }

    // Set up Informer for ConfigMaps
    factory := informers.NewSharedInformerFactory(clientset, time.Minute)
    informer := factory.Core().V1().ConfigMaps().Informer()

    // Set up WorkQueue
    queue :=
workqueue.NewRateLimitingQueue(workqueue.DefaultControllerRateLimiter())
```

```
// Event Handlers
informer.AddEventHandler(cache.ResourceEventHandlerFuncs{
    AddFunc: func(obj interface{}) {
        key, err := cache.MetaNamespaceKeyFunc(obj)
        if err == nil {
            queue.Add(key)
        }
    },
    UpdateFunc: func(oldObj, newObj interface{}) {
        key, err := cache.MetaNamespaceKeyFunc(newObj)
        if err == nil {
            queue.Add(key)
        }
    },
    DeleteFunc: func(obj interface{}) {
        key, err := cache.DeletionHandlingMetaNamespaceKeyFunc(obj)
        if err == nil {
            queue.Add(key)
        }
    },
})

// Start Informer
stopCh := make(chan struct{})
defer close(stopCh)
go factory.Start(stopCh)

// Wait for cache to sync
if !cache.WaitForCacheSync(stopCh, informer.HasSynced) {
    runtime.HandleError(fmt.Errorf("Timed out waiting for caches to sync"))
    return
}

// Process items from WorkQueue
wait.Until(func() {
    for processNextItem(queue) {
    }
}, time.Second, stopCh)

func processNextItem(queue workqueue.RateLimitingInterface) bool {
    key, quit := queue.Get()
    if quit {
        return false
    }
    defer queue.Done(key)

    // Process the item
    fmt.Printf("Processing key: %s\n", key)
    queue.Forget(key)
    return true
}
```

# A taste of GO

- ❑ the go struct for a Pod
- ❑ defining a Pod object

```
type Pod struct {  
    metav1.TypeMeta  
    metav1.ObjectMeta  
    Spec   PodSpec  
    Status PodStatus  
}
```

```
type TypeMeta struct {  
    Kind      string  
    APIVersion string  
}
```

- ❖ metav1: the API machinery package that provides metadata for API objects

```
type ObjectMeta struct {  
    Name           string  
    GenerateName   string  
    Namespace      string  
    SelfLink       string  
    UID            types.UID  
    ResourceVersion string  
    Generation     int64  
    .  
    .  
}
```

```
package main  
  
import (  
    "fmt"  
    "k8s.io/api/core/v1"  
    metav1 "k8s.io/apimachinery/pkg/apis/meta/v1"  
)  
  
func main() {  
    pod := &v1.Pod{  
        TypeMeta: metav1.TypeMeta{  
            Kind:      "Pod",  
            APIVersion: "v1",  
        },  
        ObjectMeta: metav1.ObjectMeta{  
            Name:      "example-pod",  
            Namespace: "default",  
            Labels: map[string]string{  
                "app": "example",  
            },  
        },  
        Spec: v1.PodSpec{  
            Containers: []v1.Container{  
                {  
                    Name: "example-container",  
                    Image: "nginx",  
                },  
            },  
        },  
    }  
  
    // Print the Pod object  
    fmt.Printf("Pod: %+v\n", pod)  
}
```

# The Operator Pattern

## ❑ What is an Operator?

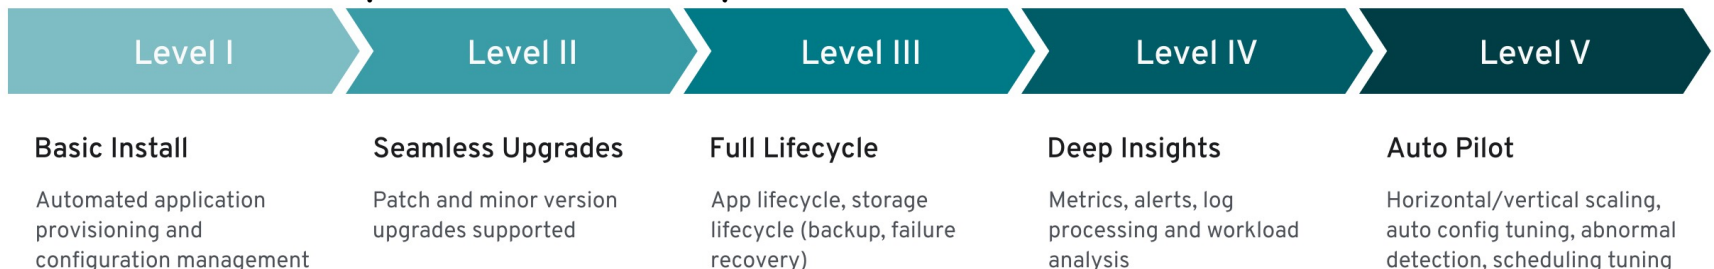
- ❖ *"An operator is a K8s controller that understands 2 domains: K8s and something else. It can automate tasks that usually require a human operator that understands both domains" - CNCF*

## ❑ CNCF Operator White Paper - Final Version

- ❖ [https://github.com/cncf/tag-app-delivery/blob/main/operator-wg/whitepaper/Operator-WhitePaper\\_v1-0.md](https://github.com/cncf/tag-app-delivery/blob/main/operator-wg/whitepaper/Operator-WhitePaper_v1-0.md)

## ❑ The registry for Kubernetes Operators

- ❖ <https://operatorhub.io/>
- ❖ 5-level operator maturity model



# Controllers vs Operators

## ❑ Controllers

- ❖ A simple reconciliation process that monitors and acts on standard K8s resources. More often, they enhance platform behavior and add new platform features.

## ❑ Operators

- ❖ A sophisticated reconciliation process that interacts with CustomResourceDefinitions (CRDs). Typically, these Operators encapsulate complex application domain logic and manage the full application lifecycle.

# The Operator Pattern - how to build?

- ❑ Use client-go library directly
- ❑ Kubebuilder
  - ❖ Owned and maintained by the K8S SIG API Machinery, a tool and set of libs
- ❑ The Operator SDK
  - ❖ From CoreOS/Red Hat
- ❑ Metacontroller, KUDO and etc.