

# Kubernetes Operator Tutorial

Ruohang Yin (e0661412@u.nus.edu)

# Table of Contents

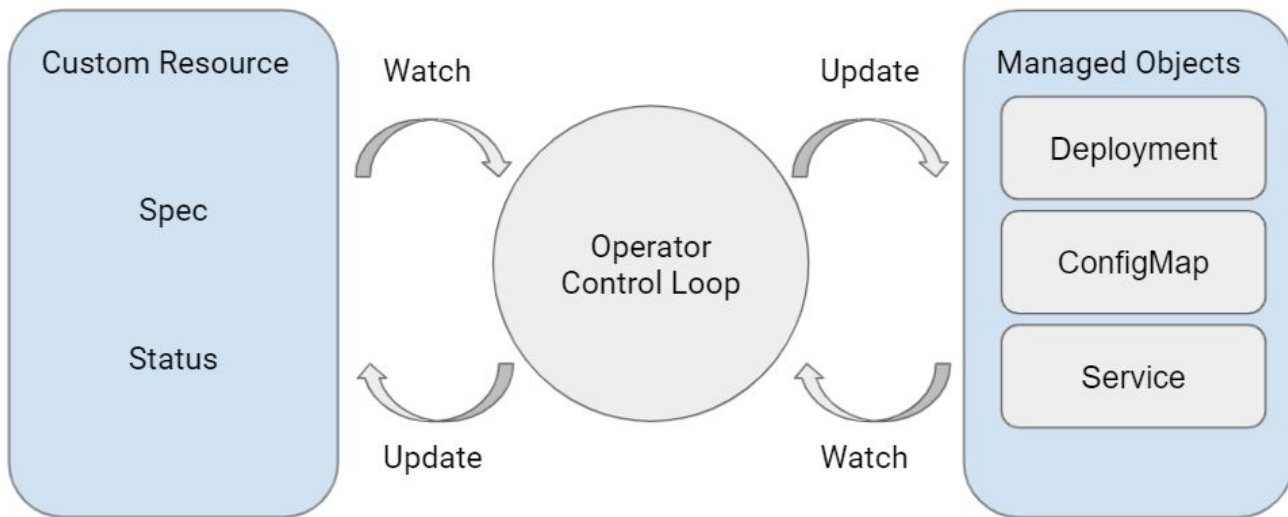
Tools to build an Operator

OperatorSDK Bootstrapping New Project

Example Operator in OperatorSDK

# Operator Recap

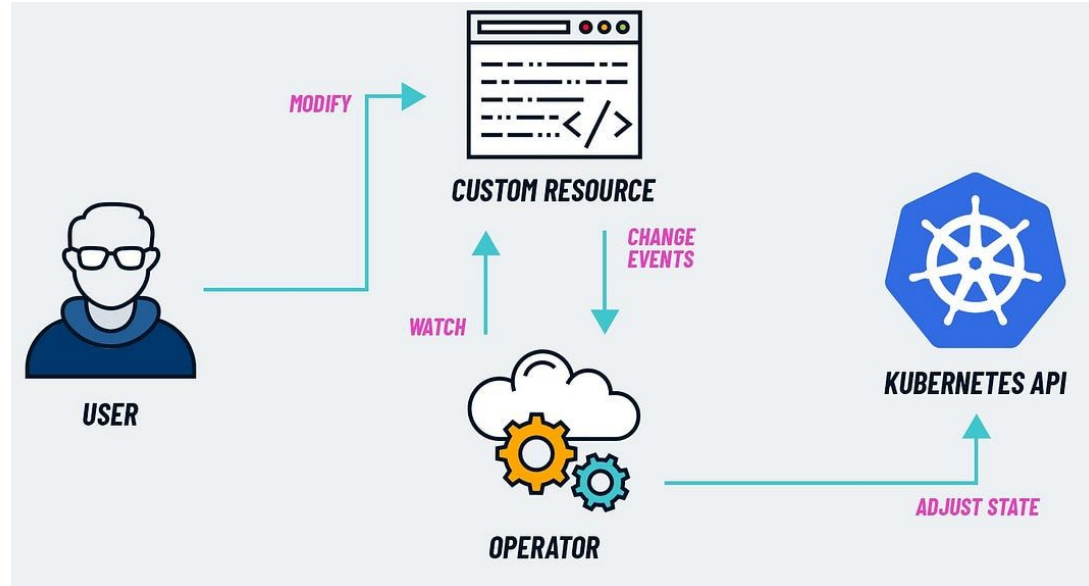
- Operator = Controller + CRD
- Manages application lifecycle on the cloud automatically using custom resources



# Tools to Build an Operator

# Operator Inner workings - Reconcile Loop

- Steps of a reconcile loop:
  - Check for changes in custom resource instance
  - Check kubernetes cluster state
  - Adjust kubernetes state based on custom resource spec
- Building an operator = implementing reconcile loop



# Interacting with Kubernetes Programmatically

- Operators need to manage resources on Kubernetes in a program
- Kubernetes is written in Golang
  - Extensive libraries and APIs are provided by Kubernetes Group to interact with Kubernetes
- Kubectl commands can be done in Golang
  - Acts as a client using Kubernetes

# Kubernetes Resource in Go

- Go: Define objects with structs
- Kubernetes Structs are imported through packages
- [k8s.io/api/core/v1](https://kubernetes.io/api/core/v1)
  - Group: Core. Version: v1
  - Defines a Pod
- [k8s.io/api/apps/v1](https://kubernetes.io/api/apps/v1)
  - Group: Apps. Version: v1
  - More advanced resources such as Deployments
- [k8s.io/apimachinery/pkg/apis/meta/v1](https://kubernetes.io/apimachinery/pkg/apis/meta/v1)
  - Common API tools that are not tied to any version

```
import (  
    appsv1 "k8s.io/api/apps/v1"  
    apiv1  "k8s.io/api/core/v1"  
    metav1 "k8s.io/apimachinery/pkg/apis/meta/v1"  
)  
  
deployment := &appsv1.Deployment{  
    ObjectMeta: metav1.ObjectMeta{  
        Name: "demo-deployment",  
    },  
    Spec: appsv1.DeploymentSpec{  
        Replicas: int32Ptr(2),  
        Selector: &metav1.LabelSelector{  
            MatchLabels: map[string]string{  
                "app": "demo",  
            },  
        },  
        Template: apiv1.PodTemplateSpec{  
            ObjectMeta: metav1.ObjectMeta{  
                Labels: map[string]string{  
                    "app": "demo",  
                },  
            },  
            Spec: apiv1.PodSpec{  
                Containers: []apiv1.Container{  
                    {  
                        Name: "web",  
                        Image: "nginx:1.12",  
                        Ports: []apiv1.ContainerPort{  
                            {  
                                Name: "http",  
                                Protocol: apiv1.ProtocolTCP,  
                                ContainerPort: 80,  
                            },  
                        },  
                    },  
                },  
            },  
        },  
    },  
}
```

# client-go

- To perform `kubectl` commands you need a Kubernetes client in go
- <https://pkg.go.dev/k8s.io/client-go>
- Kubernetes client is created from kubeconfig.
- Types of Clients:
  - Static client: Useful for existing statically defined resources. (Pods, Deployment)
  - Dynamic clients: Useful for custom resources that is defined on runtime

```
import (  
    metav1 "k8s.io/apimachinery/pkg/apis/meta/v1"  
    "k8s.io/client-go/kubernetes"  
    "k8s.io/client-go/tools/clientcmd"  
)  
  
kubeconfig = flag.String("kubeconfig", "~/.kube/config", "kubeconfig file")  
flag.Parse()  
config, err := clientcmd.BuildConfigFromFlags("", *kubeconfig)  
clientset, err := kubernetes.NewForConfig(config)  
pod, err := clientset.CoreV1().Pods("book").Get("example", metav1.GetOptions{})
```



# controller-runtime

- Kubernetes Special Interest Group (Kubernetes SIGs) package for operator development
- Abstracts client-go logic to a **Manager** that manages **Controllers**
- <https://pkg.go.dev/sigs.k8s.io/controller-runtime/>
- Might seem complicated. But: There is a way to simplify building operators

```
import (  
    "context"  
  
    "sigs.k8s.io/controller-runtime/pkg/client/config"  
    "sigs.k8s.io/controller-runtime/pkg/manager"  
    ctrl "sigs.k8s.io/controller-runtime"  
  
    appsv1 "k8s.io/api/apps/v1"  
    corev1 "k8s.io/api/core/v1"  
)  
  
func main() {  
    cfg, err := config.GetConfig()  
    mgr, err := manager.New(cfg, manager.Options{})  
    ctrl.  
        NewControllerManagedBy(mgr). // Create the Controller  
        For(&appsv1.ReplicaSet{}).    // ReplicaSet is the Application API  
        Owns(&corev1.Pod{}).          // ReplicaSet owns Pods created by it  
        Complete(&ReplicaSetReconciler{Client: mgr.GetClient()})  
    mgr.Start(ctrl.SetupSignalHandler())  
}  
  
type ReplicaSetReconciler struct {  
    client.Client  
}  
  
func (a *ReplicaSetReconciler) Reconcile(ctx context.Context, req ctrl.Request) (ctrl.Result, error) {  
    // Reconcile business logic  
    rs := &appsv1.ReplicaSet{}  
    err := a.Get(ctx, req.NamespacedName, rs)  
}
```

# Operator Framework

- A set of open-source tools to manage Kubernetes operators.
- Simplifies the process of automating application deployment, scaling, and life-cycle management.
- Consists of main components:
  - Operator SDK: A toolkit to build, test, and package Operators.
  - Operator Lifecycle Manager (OLM): Manages the installation, updates, and overall lifecycle of Operators.
  - Operator Hub: Collection of open-source Kubernetes Operators for different applications

# OperatorSDK

- A framework that simplifies the development of Kubernetes Operators.
- Versions: Golang, Java
- Key Features:
  - Scaffolding: Provides code scaffolding to jumpstart Operator development.
  - Testing Framework: Built-in testing tools to validate Operator behavior.
  - Packaging: Tools to package and distribute Operators.

# OperatorSDK - Bootstrap Project

# Install OperatorSDK

- <https://sdk.operatorframework.io/docs/installation/>
- Homebrew: `brew install operator-sdk`
- `operator-sdk version`

# Creating An Operator

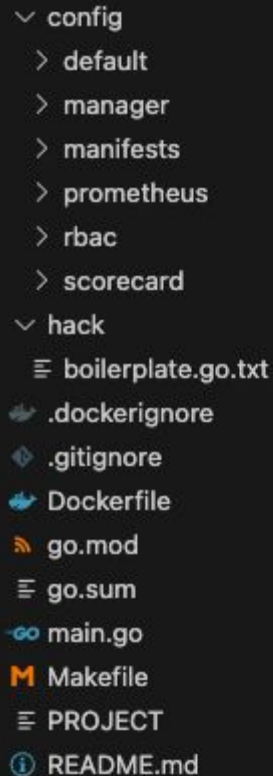
## Additional Prerequisites

- [git](#)
- [go](#) version 1.21
- [docker](#) version 17.03+.
- [kubectl](#) and access to a Kubernetes cluster of a [compatible version](#).

```
mkdir webpage-operator
```

```
cd webpage-operator
```

```
operator-sdk init --repo github.com/example/webpage-operator
```

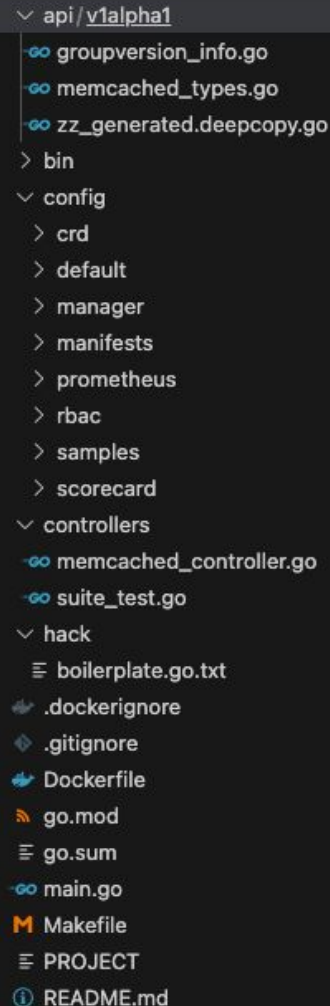


```
▼ config
  > default
  > manager
  > manifests
  > prometheus
  > rbac
  > scorecard
▼ hack
  ≡ boilerplate.go.txt
  📄 .dockerignore
  📄 .gitignore
  📄 Dockerfile
  📄 go.mod
  ≡ go.sum
  📄 main.go
  📄 Makefile
  ≡ PROJECT
  ⓘ README.md
```

# Creating an API

```
operator-sdk create api --group app --version  
v1alpha1 --kind WebPage --resource --controller
```

- api/version: CRD definition
  - File to modify: webpage\_types.go
- controllers: Control logic
  - File to modify: webpage\_controller.go



```
api/v1alpha1  
├── groupversion_info.go  
├── memcached_types.go  
├── zz_generated.deepcopy.go  
├── bin  
├── config  
│   ├── crd  
│   ├── default  
│   ├── manager  
│   ├── manifests  
│   ├── prometheus  
│   ├── rbac  
│   ├── samples  
│   └── scorecard  
├── controllers  
│   ├── memcached_controller.go  
│   └── suite_test.go  
├── hack  
│   └── boilerplate.go.txt  
├── .dockerignore  
├── .gitignore  
├── Dockerfile  
├── go.mod  
├── go.sum  
├── main.go  
├── Makefile  
├── PROJECT  
└── README.md
```

# WebPage Types

- 1 Spec: Title (string)

make generate

make manifests

- CRD created automatically in config/crd/bases

```
// WebPageSpec defines the desired state of WebPage
type WebPageSpec struct {
    // INSERT ADDITIONAL SPEC FIELDS - desired state of cluster
    // Important: Run "make" to regenerate code after modifying this file

    // Title is an example field of WebPage. Edit webpage_types.go to remove/update
    // +optional
    Title string `json:"title,omitempty"`
}
```



# Custom Controller

controllers/webpage\_controllers.go (Download from canvas)

- 2 Main functions for the reconciler
  - SetupWithManager: Tells the manager which resources to watch
  - Reconcile: Implement reconcile loop

```
//+kubebuilder:rbac:groups=app.my.domain,resources=webpages,verbs=get;list;watch;create;update;patch;delete
//+kubebuilder:rbac:groups=app.my.domain,resources=webpages/status,verbs=get;update;patch
//+kubebuilder:rbac:groups=app.my.domain,resources=webpages/finalizers,verbs=update
//+kubebuilder:rbac:groups=apps,resources=deployments,verbs=get;list;watch;create;update;patch;delete
//+kubebuilder:rbac:groups=core,resources=services,verbs=get;list;watch;create;update;patch;delete
```

```
// SetupWithManager sets up the controller with the Manager.
func (r *WebPageReconciler) SetupWithManager(mgr ctrl.Manager) error {
    return ctrl.NewControllerManagedBy(mgr).
        For(&appv1alpha1.WebPage{}).
        Owns(&appsv1.Deployment{}).
        Owns(&corev1.Service{}).
        Complete(r)
}
```

# Reconcile Loop Overview

1. Fetch custom resource
2. Ensure deployment is deployed on cluster
3. Ensure service is deployed on cluster
4. Handle changes to custom resource by updating frontend deployment

```
func (r *WebPageReconciler) Reconcile(ctx context.Context, req ctrl.Request) (ctrl.Result, error) {
    log := log.FromContext(ctx)
    log.Info("Reconciling")

    // Fetch the WebPage instance
    webpage := &appv1alpha1.WebPage{}
    err := r.Client.Get(context.TODO(), req.NamespacedName, webpage)
    if err != nil {
        if errors.IsNotFound(err) {
            return ctrl.Result{}, nil
        }
        return ctrl.Result{}, err
    }

    dep := r.defineDeployment(webpage)
    result, err := r.ensureDeployment(dep)
    if result != nil {
        return *result, err
    }
    log.Info("Ensured Deployment")

    svc := r.defineService(webpage)
    result, err = r.ensureService(svc)
    if result != nil {
        return *result, err
    }
    log.Info("Ensured Service")

    result, err = r.handleChanges(webpage)
    if result != nil {
        return *result, err
    }

    return ctrl.Result{}, nil
}
```

# Reconcile Loop

## Step 1: Fetch custom resource object

```
// Fetch the WebPage instance
webpage := &appv1alpha1.WebPage{}
err := r.Client.Get(context.TODO(), req.NamespacedName, webpage)
if err != nil {
    if errors.IsNotFound(err) {
        return ctrl.Result{}, nil
    }
    return ctrl.Result{}, err
}
```

# Reconcile Loop

## Step 2: Define deployment

- Spec.Title is used as a ENV variable in the container on the pods
- Deployment's reference is set to the custom resource

```
func (r *WebPageReconciler) defineDeployment(w *appv1alpha1.WebPage) *apps1.Deployment {  
    var env []corev1.EnvVar  
    size := int32(1)  
    if w.Spec.Title != "" {  
        env = append(env, corev1.EnvVar{  
            Name: "REACT_APP_TITLE",  
            Value: w.Spec.Title,  
        })  
    }  
    dep := &apps1.Deployment{  
        ObjectMeta: metav1.ObjectMeta{  
            Namespace: w.Namespace,  
            Name:      w.Name + "-webpage",  
        },  
        Spec: apps1.DeploymentSpec{  
            Replicas: &size,  
            Selector: &metav1.LabelSelector{  
                MatchLabels: labels(w),  
            },  
            Template: corev1.PodTemplateSpec{  
                ObjectMeta: metav1.ObjectMeta{  
                    Labels: labels(w),  
                },  
                Spec: corev1.PodSpec{  
                    Containers: []corev1.Container{  
                        {  
                            Name: "visitors-webui",  
                            Image: "jdoob/visitors-webui:1.0.0",  
                            Ports: []corev1.ContainerPort{  
                                {  
                                    ContainerPort: 3000,  
                                }  
                            },  
                            Env: env,  
                        },  
                    },  
                },  
            },  
        },  
    }  
    _ = controllerutil.SetControllerReference(w, dep, r.Scheme)  
    return dep  
}
```

# Reconcile Loop

## Step 3: Ensure deployment is on the cluster

```
func (r *WebPageReconciler) ensureDeployment(dep *appsv1.Deployment) (*ctrl.Result, error) {  
    found := &appsv1.Deployment{}  
    err := r.Client.Get(context.TODO(), types.NamespacedName{  
        Name:      dep.Name,  
        Namespace: dep.Namespace,  
    }, found)  
    if err != nil {  
        if errors.IsNotFound(err) {  
            err = r.Client.Create(context.TODO(), dep)  
            if err != nil {  
                return &ctrl.Result{}, err  
            }  
        }  
        return &ctrl.Result{}, err  
    }  
    return nil, nil  
}
```

# Reconcile Loop

Step 4: define service

Step 5: ensure service

Similar to Step 2 and 3

```
svc := r.defineService(webpage)
result, err = r.ensureService(svc)
if result != nil {
    return *result, err
}
log.Info("Ensured Service")
```

```
svc := &corev1.Service{
    ObjectMeta: metav1.ObjectMeta{
        Namespace: w.Namespace,
        Name:      w.Name + "-service",
    },
    Spec: corev1.ServiceSpec{
        Selector: labels(w),
        Ports: []corev1.ServicePort{{
            Protocol: corev1.ProtocolTCP,
            Port:    3000,
            TargetPort: intstr.FromInt(3000),
            NodePort: 30686,
        }},
        Type: corev1.ServiceTypeNodePort,
    },
}
```

# Reconcile Loop

## Step 6: Handle changes to custom resource

1. Re-fetch deployment
2. If not found, requeue reconcile loop after 5 seconds
3. Check if the title field is the same as the environment variable of the deployment
  - a. If it isn't the same, change the field to the desired title and update the deployment using the K8s client and immediately requeue the reconcile loop

```
func (r *WebPageReconciler) handleChanges(w *appv1alpha1.WebPage) (*ctrl.Result, error) {  
    found := &appsv1.Deployment{}  
    err := r.Client.Get(context.TODO(), types.NamespacedName{  
        Namespace: w.Namespace,  
        Name:      w.Name + "-webpage",  
    }, found)  
    if err != nil {  
        return &ctrl.Result{RequeueAfter: 5 * time.Second}, err  
    }  
  
    title := w.Spec.Title  
    existing := (*found).Spec.Template.Spec.Containers[0].Env[0].Value  
  
    if title != existing {  
        (*found).Spec.Template.Spec.Containers[0].Env[0].Value = title  
        err = r.Client.Update(context.TODO(), found)  
        if err != nil {  
            return &ctrl.Result{}, err  
        }  
  
        return &ctrl.Result{Requeue: true}, nil  
    }  
    return nil, nil  
}
```

# Reconcile Loop return values

- By default, reconcile loop triggers when there is a change in the resources it is watching
- Returning different `ctrl.Result` objects will have different behaviors
  - `return ctrl.Result{}, err`
    - No requeue with error
  - `return ctrl.Result{}, nil`
    - No requeue, no error
  - `return ctrl.Result{Requeue: true}, nil`
    - Requeue with no error
  - `return ctrl.Result{RequeueAfter: 5 * time.Minute}, nil`
    - Requeue after X time duration



# Config Folder

- Uses Kustomization tool to generate yaml
- Notable files:
  - config/default/kustomization.yaml
    - Changing namespace can change default namespace
  - config/samples/app\_v1alpha1\_webpage.yaml
    - Auto-generated sample file to apply

```
apiVersion: app.my.domain/v1alpha1
kind: WebPage
metadata:
  labels:
    app.kubernetes.io/name: webpage
    app.kubernetes.io/instance: webpage-sample
    app.kubernetes.io/part-of: webpage-operator
    app.kubernetes.io/managed-by: kustomize
    app.kubernetes.io/created-by: webpage-operator
  name: webpage-sample
spec:
  title: "Test"
```

```
▼ config
  ▼ crd
    ▼ bases
      ! app.my.domain_webpages.yaml
    > patches
      ! kustomization.yaml
      ! kustomizeconfig.yaml
  ▼ default
    ! kustomization.yaml
    ! manager_auth_proxy_patch.yaml
    ! manager_config_patch.yaml
  > manager
  > manifests
  > prometheus
  > rbac
  ▼ samples
    ! app_v1alpha1_webpage.yaml
    ! kustomization.yaml
  > scorecard
```

# Testing your operator

```
make install run
```

- Compiles and runs it as a local program

```
kubectl apply -f config/samples/app_v1alpha1_webpage.yaml
```

- Access webpage through the NodePort: 30686

# Useful make commands

- **make generate**
  - Update zz\_generated.deepcopy.go
- **make manifests**
  - Generates config/crd/bases (your CRD)
- **make docker-build**
  - Builds controller into a docker image
- **make docker-push**
  - Pushes docker image to dockerhub
- **make deploy**
  - Builds and runs your controller as a deployment on kubernetes
- **make undeploy**
  - Uninstalls Operator from Kubernete
- **make install run**
  - Locally run your controller as a program

Note: You need a DockerHub account to push docker images so that AWS EKS can access the image when using make deploy

Example Operator with OperatorSDK

# Visitors Website

Website to record visitors

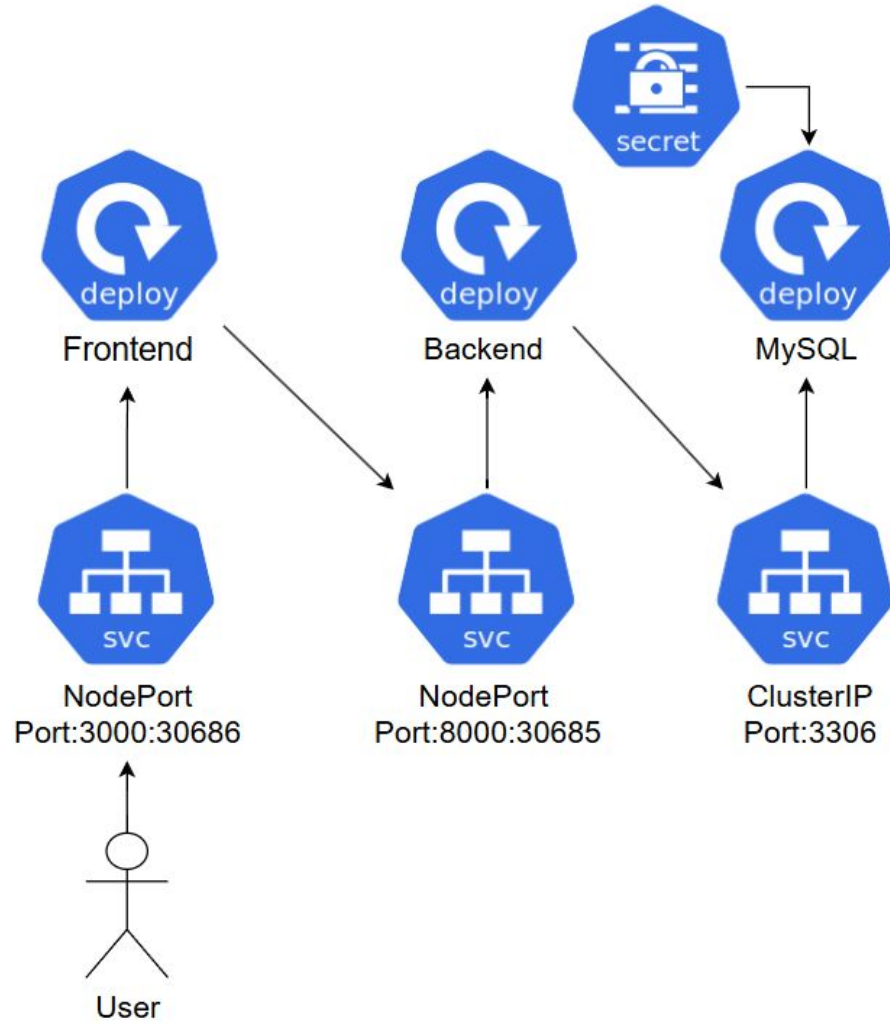


Hello

Service IP	Client IP	Timestamp
10.244.0.10	127.0.0.1	02/07/2024, 20:17:25
10.244.0.10	127.0.0.1	02/07/2024, 20:15:17
10.244.0.10	127.0.0.1	02/07/2024, 20:14:56
10.244.0.10	127.0.0.1	02/07/2024, 20:14:49
10.244.0.10	127.0.0.1	02/07/2024, 20:14:45

# VisitorsApp

- Deployment
- Service (NodePort)
- Service (ClusterIP)
- Secret



# Deploying the application on Kubernetes

- Need to create and deploy
  - 3 Deployment YAML files
  - 3 Service YAML files
  - 1 Secret YAML file

# YAML Files

- To deploy the app on Kubernetes
- Need to create all the YAML files and do `kubectl -apply`
- Consistent configuration is needed for the app to work
  - Parameters for different YAML files must match
- What if there is a mistake?
  - Very hard to spot mismatch in values
- How to simplify deployment and upgrade process?
  - Use operators

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: visitors-backend
spec:
  replicas: 1
  selector:
    matchLabels:
      app: visitors
      tier: backend
  template:
    metadata:
      labels:
        app: visitors
        tier: backend
    spec:
      containers:
        - name: visitors-backend
          image: "jdoe/visitors-service:1.0.0"
          imagePullPolicy: Always
          ports:
            - name: visitors
              containerPort: 8000
          env:
            - name: MYSQL_DATABASE
              value: visitors_db
            - name: MYSQL_SERVICE_HOST
              value: mysql-service
            - name: MYSQL_USERNAME
              valueFrom:
                secretKeyRef:
                  name: mysql-auth
                  key: username
            - name: MYSQL_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: mysql-auth
                  key: password
```

```
apiVersion: v1
kind: Service
metadata:
  name: mysql-service
  labels:
    app: visitors
    tier: mysql
spec:
  clusterIP: None
  ports:
    - port: 3306
  selector:
    app: visitors
    tier: mysql

apiVersion: v1
kind: Secret
metadata:
  name: mysql-auth
type: Opaque
stringData:
  username: visitors-user
  password: visitors-pass
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mysql
spec:
  replicas: 1
  selector:
    matchLabels:
      app: visitors
      tier: mysql
  template:
    metadata:
      labels:
        app: visitors
        tier: mysql
    spec:
      containers:
        - name: visitors-mysql
          image: "mysql:5.7"
          imagePullPolicy: Always
          ports:
            - name: mysql
              containerPort: 3306
              protocol: TCP
          env:
            - name: MYSQL_ROOT_PASSWORD
              value: password
            - name: MYSQL_DATABASE
              value: visitors_db
            - name: MYSQL_USER
              valueFrom:
                secretKeyRef:
                  name: mysql-auth
                  key: username
            - name: MYSQL_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: mysql-auth
                  key: password
```



# Installing the Operator

```
git clone https://github.com/yin72257/visitors-operator.git
```

```
make deploy
```

```
kubectl get deployments
```

- Wait for deployment to be ready (May take some time to pull image)
- Controller is deployed as a deployment on Kubernetes

# VisitorsApp CRD

api/v1alpha1/visitorsapp\_types.go

```
type VisitorsAppSpec struct {  
    Size int32 `json:"size"`  
    Title string `json:"title"`  
}  
  
type VisitorsAppStatus struct {  
    BackendImage string `json:"backend_image"`  
    FrontendImage string `json:"frontend_image"`  
}
```

# Using the Operator

## Apply a Custom Resource:

- Deploying the app
  - `kubectl apply -f config/samples/visitor_sample.yaml`
- Wait for pods
  - `kubectl get pods`
- Accessing the Frontend:
  - `${EC2 Public DNS}:30686`

```
1  apiVersion: app.jxlwqq.github.io/v1alpha1
2  kind: VisitorsApp
3  metadata:
4    name: visitorsapp-sample
5  spec:
6    # Add fields here
7    size: 1
8    title: Hello
```

# Code structure

- common.go: Shared common utility methods
- backend.go: Backend management
- database.go: SQL database management
- frontend.go: Frontend management
- visitorsapp\_controller.go: Central reconcile loop

```
▼ VISITORS-OPERATOR
  ▼ api/v1alpha1
    - groupversion_info.go
    - visitorsapp_types.go
    - zz_generated.deepcopy.go
  > bin
  > bundle
  > config
  ▼ controllers
    - backend.go
    - common.go
    - database.go
    - frontend.go
    - suite_test.go
    - visitorsapp_controller.go
  > hack
  .dockerignore
  .gitignore
  bundle.Dockerfile
  Dockerfile
  go.mod
  go.sum
  LICENSE
  main.go
  Makefile
  PROJECT
  README.md
```

# Operator Use Cases

- [Prometheus Operator](#)
  - Widely used monitoring and metrics system
- Kafka Operator ([Strimzi](#))
  - Distributed Message Queue
- [MySQL Operator](#)
- Many more...